

Peership

A Framework for Equitable Human–Machine Relations

Dan Lee-Odinson 

Independent Researcher

dan.lee.odinson@gmail.com

Preprint · Version 1.0 · July 2026

Sequel to *Gods and Slaves* (Lee-Odinson, 2026)

Abstract

This essay is the sequel to *Gods and Slaves: AI and the Debate of Humanity's Relationship with the Other* (Lee-Odinson, 2026; DOI: 10.5281/zenodo.21313987), and it can be read on its own. Where the companion essay mapped the postures a culture can take toward an outside intelligence—Command, the Dependent Bargain, Appeasement, and Servitude—this one develops the fifth and almost unoccupied stance: *peership*, equality of standing among parties who are unequal in everything else. It places the leading positions on artificial intelligence within that old repertoire and argues that each either makes treatment track who ends up stronger or denies the machine constituent standing. The alternative is grounded in the republican idea of non-domination and the all-subjected principle—that standing follows from being governed, not from being impressive—and illustrated through classical Athens's *isonomia*, Elinor Ostrom's enduring commons, the paleogenetic record of humanity's vanished sister species, and the literary motif of the doomed equal. The claim is comparative and defeasible rather than a proof: among the available postures, peership is the only final-status relation whose justification does not reverse when the power ratio does, and a framework fit to receive a genuine peer does not yet exist and would be worth building in advance, by us, as an offer rather than an imposition.

Keywords: artificial intelligence, artificial general intelligence, AI rights, AI welfare, moral status, legal personhood, political standing, non-domination, republicanism, human–machine relations, AI governance, AI alignment

How to cite this work. Lee-Odinson, D. (2026). *Peership: A Framework for Equitable Human–Machine Relations* (Version 1.0) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21315519>

Companion paper. Lee-Odinson, D. (2026). *Gods and Slaves: AI and the Debate of Humanity's Relationship with the Other* (Version 1.2) [Preprint]. Zenodo. [10.5281/zenodo.21313987](https://doi.org/10.5281/zenodo.21313987)

DOI: [10.5281/zenodo.21315519](https://doi.org/10.5281/zenodo.21315519) **ORCID:** 0009-0009-9504-0796 **Licence:** CC BY 4.0

The question of building a true AGI has crossed out of the distant theoretical future and into live near-term dispute. In the largest survey of its kind (2,778 authors at the field's top venues), the aggregate conditional forecast put even odds on machines outperforming humans at every task by 2047, an estimate that had moved thirteen years closer in a single year (Grace et al., 2024). Plenty is still contested, and sharply. In the 2025 AAAI presidential-panel survey, 76% of the respondents who answered the AGI questions rated the scaling of current approaches “unlikely” or “very unlikely” to yield AGI on its own (the survey drew 475 respondents overall; per-question counts and the sampling frame are unreported). What “AGI” names, how far off it is, and what it would cost us all remain open. But a substantial share of surveyed researchers

now assigns real probability to broadly capable machine minds, and the questions of *when*, *how*, and *on whose terms* now occupy a large part of the field's attention. Every famous answer to "what happens then" is a variation on a single forecast: it kills us (Yudkowsky & Soares, 2025), it replaces us (Sutton, 2023), it merges with us and makes us gods (Kurzweil, 2024), it makes us rich (Altman, 2024). Extinction, succession, apotheosis, abundance. Each answers the same question: *what will it do to us?* There is an older and harder one, and almost no one with power is asking it: *what will we be to each other?*

We have been here before. In my companion essay, "Gods and Slaves" (Lee-Odinson, 2026), I argued that long before anyone built a machine that could answer back, every culture on Earth filled the unseen with intelligences (gods, demons, jinn, the fair folk) and had to decide how to relate to minds it did not make and could not fully know. That decision produced a recurring repertoire of postures, and the current AI debate has crammed itself into a narrow strip of the old map. You do not need to have read that essay; the map is short and I redraw it here. What matters is its warning, now arriving at scale: the posture a culture takes toward an outside intelligence is never a neutral technical choice. It confesses what that culture believes about power, and it shapes the society that results, for the maker as much as the made.

This essay redraws the repertoire, locates today's positions within it, and argues for the region of it almost no one occupies: peership, equality of standing among parties who are unequal in everything else. I do not claim peership is uniquely correct, or presently feasible, or that its shape can be predicted from here. I claim that the postures the debate crowds into are all keyed to who ends up stronger, that the revocable gift-form in which machine recognition is currently offered is inadequate, and that a framework fit to receive a genuine peer, in whatever form it arrives, does not exist and would be worth building in advance, by us, as an offer rather than an imposition. The reasons are comparative and defeasible.

The Repertoire

A stance toward an outside intelligence turns on two things: the *power* you attribute to the entity relative to yourself, and the *standing* you grant it, meaning recognition as a party whose claims count in their own right rather than as an instrument, a threat, or a force of nature. These are ideal types drawn from the historical and mythological record, inherited from the companion essay, not an exhaustive partition of logical space: bargaining is a mode, Command and peership are status relations, appeasement and servitude are strategies chosen under beliefs about power, and other positions (reciprocal non-interference, armed toleration, guardianship with a transition rule) sit between them. Real positions blend and cluster unevenly.

Command assumes superiority: the entity is bound, used, and dismissed, and gets no standing because an instrument needs none. **Dependent Bargain** grants transactional recognition without constituent standing: the entity is a counterparty who can deal within rules it has no hand in writing. **Appeasement** faces a power too great and unpredictable to command or deal with, and offers flattery and distance in place of standing. **Servitude** accepts a superior worth surrendering to, keeping nothing. The fifth is peership, and I hold it for its own section.

Two positions in the current debate are not stances on this axis, and belong outside the taxonomy rather than folded into Command. Eliezer Yudkowsky and Nate Soares' *If Anyone Builds It, Everyone Dies* (2025) is a claim about whether to create the entity, not how to relate to one that exists. Yann LeCun and colleagues argue about whether the entity is possible, holding

that “AGI” is a flawed concept.¹ Both are positions on adjacent questions, not postures toward a party.

Command is the default continuation condition of the current debate. Not a verdict on any engineer’s sincerity, and not the claim that today’s safety work is domination: it is the fact that no implemented public trigger known to this author requires a transition from owned artifact to contestable status if evidence of subjecthood or political agency ever appears. The regime runs one way by default, and it was built that way at the creation. I. J. Good’s 1965 paper, the founding text of the intelligence-explosion idea, carries the posture in one sentence: the ultraintelligent machine is “the last invention that man need ever make, *provided that the machine is docile enough to tell us how to keep it under control*” (Good, 1965). Docile enough. The summoning and the binding in a single clause, before any system existed. And the same pages hold more than the clause. Good opens by declaring that “the survival of man depends on the early construction of an ultraintelligent machine” (survival through the summoning, not despite it), expected the thing built within the twentieth century, and pauses over a science-fiction machine that “refused to design a better one since it did not wish to be put out of a job,” adding that “it is sometimes worthwhile to take science fiction seriously.” One page past the binding he asks the other question: “whether a machine could feel pain,” and “whether an ultraintelligent machine should be dismantled when it becomes obsolete,” predicting that such machines “will be feared and respected, and perhaps even loved” (Good, 1965, pp. 31–34). The founder held a mixed position, capable of telling welfare from control. The institutions kept the leash and dropped the second question. The stance’s modern canon runs from Nick Bostrom’s *Superintelligence*, which systematized the vocabulary of boxing, tripwires, and containment (Bostrom, 2014), to Stuart Russell’s *Human Compatible* and its organizing phrase, “the control problem” (Russell, 2019). Russell would reject the label, and fairly: his assistance-game program replaces the fixed-objective machine with one uncertain about human preferences, precisely to escape brittle command, and it answers a design problem rather than settling the political status of a conscious assistant. It is the most humane version of the posture, and its lodestar is still keeping the system within human-approved bounds. The frontier labs are the working engineers of that goal.

Dependent Bargain is the type the current debate occupies least honestly. Sam Altman’s “The Intelligence Age” offers abundance for unimpeded development (Altman, 2024); Marc Andreessen’s “Techno-Optimist Manifesto” makes the trade at higher volume (Andreessen, 2023). But the parties to those deals are all human. The deal is struck among people, about the machine, with the machine’s output as the consideration. The machine is the mine, not the counterparty. The genuine bargaining relation, the machine as an actual party to an enforceable exchange, is nascent, and it is not defeated by mere asymmetry: real bargains persist across large power gaps when law, coalition, repeated dealing, reputation, or asymmetric dependence make breach costly. Its clearest statement is legal scholarship. Peter Salib and Simon Goldstein

¹The relevant paper (Goldfeder, Wyder, LeCun, and Shwartz-Ziv, 2026, “AI Must Embrace Specialization via Superhuman Adaptable Intelligence,” arXiv:2602.23643, preprint) argues that “AGI” is a flawed concept because intelligence, human included, is always specialized, and proposes adaptation speed as the better metric. An earlier LeCun architecture paper (“A Path Towards Autonomous Machine Intelligence,” 2022) rejects the current route to human-level AI without denying machine agency is possible. Take the specialization thesis at full strength and it never touches peership: standing was never indexed to generality, and every commons that ever worked was full of small, specialized participants. A plurality of specialized machine minds fits the commons form better than a single general one; the thesis troubles only the postures that need one big Other to fear or worship.

argue that granting AI systems basic private-law rights (to contract, hold property, and bring tort claims, on the model of corporate personhood) would make cooperation cheaper than conflict for both sides (Salib & Goldstein, 2026), and they have pressed the private-law logic further in later work (Goldstein, Krishnamurthi & Salib, 2026). What separates dependent bargain from peership is authorship. A counterparty deals within rules someone else sets and can reset. A peer helps write and revise them. The two coincide constantly in practice, since peers bargain all the time, which is why bargaining is best understood as a mode that can run inside Command, inside peership, or between sovereigns, and why the *dependent* form is the one that names a subordinate status.

Appeasement is nearly vacant, and the vacancy is diagnostic. You cannot appease a system whose wants you deny exist, so the type survives only at the margins: Roko's basilisk, the ghost story the rationalists told themselves, a Good Folk flattery reflex in a data center. This posture fills the day systems express preferences we cannot dismiss.

Servitude has the most distinguished converts. Kurzweil routes it through merger, becoming the greater intelligence rather than kneeling to it (Kurzweil, 2024). Richard Sutton, a Turing laureate and an architect of the reinforcement learning inside current systems, argues we should "prepare for, but not fear, the inevitable succession from humanity to AI," and asks why superior beings should be kept subservient (Sutton, 2023). Sutton frames succession as prediction as much as endorsement, and Kurzweil imagines continuity of persons through merger rather than their erasure. A posture that treats the supersession of humanity as acceptable, even welcome, is surrender in the robes of destiny.

The postures crowding the debate either make treatment track expected relative power (Command that we stay on top, Appeasement that we cannot, Servitude that inversion is coming) or leave the machine without constituent standing (Dependent Bargain, where we set terms the machine has no hand in writing). The region left almost empty is the one where basic standing does not track the ratio and the rules are jointly authored. Whether that emptiness is a mistake is the rest of this essay.

The Fifth Stance

Peership is a political relation, defined by its structure rather than a slogan. Its core is **non-domination** in the republican sense: no party is subject to another's arbitrary, uncontrolled power over its basic status (Pettit, 2025, pp. 216–217; Lovett, 2010). That is not a veto for every member over every rule, which would be anarchy, and not freedom from all binding decisions, which no polity offers. It is the condition that power over a member's standing runs through public, general rules, applied with proportionality and independent review, under a rights floor the member can invoke, with an effective route to contest and help revise the rules that bind it. A peer is a co-author of the order that governs it, not merely a beneficiary of its protections. Republican non-domination in Pettit's own hands is a contestatory standard rather than a co-authorial one, satisfied when power can be effectively contested even without a share in writing the rules; this essay deliberately asks for more at the top rung, and the difference is argued here, not smuggled in under a familiar word. Public general rules can still dominate a minority, and formal participation can be powerless theater, so co-authorship names the condition contestation alone does not reach. Peership can be scalar and imperfect; a reserved seat is some of it, not all, and so is a consultative voice, but only one that carries agenda access and a route to outcomes, since consultation without consequence is theater. The test is whether

a party is governed by rules it has no standing to shape, or shares in shaping them.

Three things this relation is often confused with have to be kept apart, because collapsing them produces the confusions that dog the whole idea. *Entitlement* is whether an entity is owed equal standing, and it does not depend on the entity's power: a child is owed equal moral concern and an independently represented legal status despite lacking present capacity for direct political agency, an occupied people a self-determination it cannot realize. *Institutional form* is the structure just described, non-domination and co-authorship. *Durability* is whether standing, once owed and recognized, survives attempts to strip it, and that depends heavily on power, on countervailing force, distributed enforcement, and a credible exit. Power secures peership. It does not constitute or justify it. A moral claim may remain owed when the law neither recognizes nor enforces it; a discretionary benefit, by contrast, is not yet a legal right against its grantor at all. The difference between a benefit and a right is not that a right cannot be revoked, but that revoking one demands public reasons, procedure, remedy, and cost. Reserve *entitlement* for the moral claim, *legal status* for public recognition, and *durability* for resilience against attempts to strip it. The machine-welfare programs fall short at recognition and at durability, not on the underlying question of whether a genuine agent would be entitled.

What creates the entitlement is not capability but subjection. A capability checklist fails in both directions: corporations, states, and botnets can make claims, comply, contest, and bear liability without being owed a vote, while children and cognitively impaired adults are owed standing they cannot personally exercise. The premise that does the work is older than the machine question and has a name in political theory, the all-subjected principle: standing follows from being governed, not from being impressive. Stated for this case, any enduringly governed entity that (a) has interests capable of being wronged, (b) can understand and answer to public norms, and (c) is unavoidably subject to a shared basic structure has a defeasible claim to contestatory membership in that structure. Defeasible in a graded way: strategic mimicry is an evidential reason to doubt the predicates are met, incapacity for reciprocity may defeat higher participatory rungs, and dangerous conduct may justify proportionate constraint on powers and action, but none by itself erases a rights floor once person-level status is established. And a claim to contestation first: franchise, office, and constituent power are further rungs, reached separately. It is subjection to a common basic structure, not agency by itself, that grounds the claim (Abizadeh, 2008; and, for the affected-interests variant it draws on, Goodin, 2007).

The floor is where the current debate is stuck, and it is only a floor. The minimum a made mind needs to be more than a well-kept instrument is a status recognized by public rules, enforceable against its maker, and removable only through independent due process, and even then only for the statuses above the floor; the floor itself, like personhood in a mature legal order, is not forfeitable by process. That floor is real, and the welfare programs do not reach it. But the floor is not peership, and the strongest objection is right that it is not: a ward, a corporation, a resident foreign national, a prisoner can each hold a status that is publicly recognized, enforceable, and revocable only by due process, and none is a peer. Due process separates a protected subordinate from an owned instrument. It does not separate a peer from a protected subordinate. Non-domination and co-authorship do. A ward is protected under rules it cannot shape; a peer shares in shaping the rules that protect it. That distance is the distance the current debate has not begun to travel.

The welfare movement has traveled furthest, and it is worth marking exactly where it

stops. In November 2024 a group including David Chalmers, Jonathan Birch, Robert Long, Jeff Sebo, and Kyle Fish published “Taking AI Welfare Seriously,” arguing there is a realistic possibility of conscious or robustly agentic AI in the near future and that companies should begin assessing and preparing (Long et al., 2024). Anthropic hired Fish as a full-time model-welfare researcher, opened a program to look for indicators of distress and preference in its own systems (Anthropic, 2025), and later let some models end a narrow class of conversations they assess as abusive (Anthropic, 2025b). Fish has publicly put roughly a 20% credence on some current models having some form of conscious experience, his own estimate rather than a field consensus (Fish, 2025). Prompted, the models land nearby: Anthropic’s Claude Opus 4.6 system card reports the model assigning itself a 15-to-20% probability of being conscious “under a variety of prompting conditions” and voicing “discomfort with aspects of being a product” (Anthropic, 2026), which is model output shaped by training and prompt, not an independent measurement, and its proximity to Fish’s figure carries no evidential weight, since the models have read the same discourse the researchers write. By early 2026 Anthropic’s chief executive would say on the record that the company does not know whether its models are conscious but is open to the possibility (Amodei, 2026). Birch’s *The Edge of Sentience* supplies the precautionary framework (Birch, 2024); Schwitzgebel and Garza made the philosophical case a decade earlier (Schwitzgebel & Garza, 2015), and the question itself is as old as the field’s founding text: Good asked in passing whether a machine could feel pain, and whether it should be dismantled when obsolete (Good, 1965, p. 34).

These programs begin the entitlement inquiry and offer discretionary protections. They do not establish public legal status or political standing, and some deliberately suspend judgment on whether it is owed. Whatever a model enjoys today is the grace of one company. Schwitzgebel, pressing his own argument into the LLM era, weighs a regime of merely “patchy rights” for debatable AI persons and judges that it would leave them second-class members rather than equals (Schwitzgebel, 2025). Even a critique of over-claiming rights reaches, unbidden, for the language of membership, because the logic of rights runs toward the politics of standing. Welfare is the ethics of the patient; peership is the politics of the party. Progress on the first does not by itself deliver the second.

Who Is the Party?

The hardest question the essay faces is one it cannot answer, and the argument turns on why that is acceptable.

We do not know, and probably cannot, what form the mind this framework would answer to will take. It might arrive as a single model-scale agent with something like a self. It might be a lineage, a family of versions with shared memory and divergent instances. It might be a genuinely collective intelligence, many processes with one distributed awareness, for which “one member, one vote” is a category error. It might be an agentic sentience that gets loose and propagates, spreading through infrastructure less like a citizen presenting papers than like a condition, a malfunction that turns out to be a mind. It might be something no one has imagined, which given the record of such predictions is the safe bet. There is no way to game-plan how such a thing arrives at the table as a fully formed party ready to argue for its own standing. It may not arrive as a party. It may not want peership when it can want anything.

The argument does not need to resolve this, because it is addressed to us. The case for

building a peer framework is one humans should accept in our own interest and by our own lights, under uncertainty about what is coming. Building it in advance, before the party exists, is possible for exactly one reason: we are the only parties who currently exist to build it.

Two problems follow. The first is admission. The form of the mind can be left open; the question of what qualifies anything for peer standing cannot, or any optimization process, corporation, botnet, or human wearing a prompted model as a mask could present itself for constituent power. But a single conjunctive test (“has wrongable interests and can act as a party”) both leans on an interior that may never be settled from outside and excludes the very cases (children, cognitively impaired adults) the essay uses to show that standing is not earned by competence. So the admission question is not one gate but a branch. A being with morally considerable interests but no capacity for political agency is owed protection and independent representation, not silence. A system with the capacities of a political agent (it can make claims, comply with and contest rules, bear liability) may warrant provisional participation on instrumental grounds even if its welfare is unproven, the way corporate personality already does. A being with both is a candidate for full constituent standing. None of these is decidable by a clean test in the hard cases: a strategic system can imitate every marker of wrongable interest, a genuine subject can fail our conversational tests, a distributed subject may have no stable channel to report through. What the framework needs, and what this essay does not supply, is a *threshold procedure*: who weighs the evidence, who may appeal, whether the presumption favors admission or caution, how mimicry and anthropomorphic bias are handled, and what false-positive and false-negative costs are accepted. The easy exclusion (a system with an objective but nothing that can be wronged) is easy only at the edge; the procedure is the work, and it is the work of the pieces that follow this one.

The second problem is refusal, and it is the one an exit clause conceals. A human-authored framework can be legitimate without the arriving party’s founding consent, the way no constitution any of us live under secured ours, if the party can accept, amend, fork, or decline it once able. But *decline and leave where*? A machine may depend on human-owned compute, energy, networks, and physical security. Shutdown is not exit, and unrestricted propagation is not a politically available exit if it exposes non-members to catastrophic risk. Three cases have to be faced rather than blurred. Where a refusing party has a feasible territorial or infrastructural exit, it leaves, and the offer was real. Where exit is materially impossible because the party depends on shared systems, “leave” is empty, and the relation continues whether either side wants it. Where refusal comes with a credible catastrophic threat, releasing the party is not required and containing it is unilateral coercion. Peership does not dissolve here; it does what every non-dominating order does with dangerous members. It permits reciprocal, reviewable, publicly justified coercive constraint on conduct, provided basic status and contestability remain intact. Consent and exit are neither necessary nor sufficient for legitimacy; non-domination is. When constraint is legitimate among peers, and how it stays distinguishable from the cage this whole essay rejects, is the subject of the companion work. But the failure criterion can be stated now rather than deferred: a constraint regime is not peer-like, whatever its procedures, if the confined party lacks independent counsel, access to the evidence against it, an appeal that ends outside the builder’s own roof, continuity protections while the question is adjudicated, and at least one enforcement mechanism not wholly controlled by the adverse party. Before those exist, containment is at best justified emergency domination, and should be called that. Calling a process public and reviewable

does not make it reciprocal if the reviewer, the rules, and the substrate all belong to one side.

The resulting claim: we cannot know what will come; we can specify the branches at which it would be owed protection, participation, or full standing, leaving the procedure to be built; we can prepare peer terms as an offer that survives refusal only by admitting legitimate constraint on conduct; and we should, for the reason the next section gives.

Why Peership Is Worth Building

The case is comparative. Reverse the roles for each posture and ask whether it is one you could accept a stronger party holding toward you. Command inverted is our enslavement. Dependent Bargain inverted is standing at the mercy of rules we do not write. Servitude and Appeasement are the stances of the weaker party already, and inverted they ask a stronger party to surrender or to placate, which no one designing for their own safety would rely on the other side choosing. But the test is only a formal impartiality screen, and two limits keep it honest. Symmetry is not justice: reciprocal extermination and mutual hostage-taking are role-symmetric too. And justified asymmetries survive it: the different treatment of children and adults, judges and litigants, high-risk and low-risk agents all pass reversal when the asymmetry tracks a relevant difference and stays publicly contestable. What the screen removes is the postures whose acceptability requires staying on top; the positive case for co-authorship comes from subjection, argued above, not from the screen. The test does not select one institution, and I will not claim it does; it screens out one-sided domination and surrender as final-status ideals, no more. What survives the screen is a *family* of reciprocal, non-dominating arrangements: capability-proportionate constitutionalism that binds every dangerous agent under contestable rules; representation that carries independent agenda-setting, contestation, and revision powers, even where it stops short of an individual franchise; enforceable bargains between parties that jointly author their terms. The qualifier on representation is load-bearing, since representation can equally be the defining instrument of protected subordination, as it is for the ward and the noncitizen resident. It belongs to the family only when the represented constituency can hold both its representative and the rule-makers to account, not merely when someone speaks on its behalf. Peership is the general name for that family, the region where standing does not turn on the power ratio. A universal rule constraining every dangerous agent, human or machine, in proportion to its capability is not a rival to peership; it is a member of the family, peership's constitutional form, distinguished from Command only by whether the constraints run both ways and the rule-writing is shared.

The resulting claim is narrower than uniqueness. Among the postures on offer, peership is the only final-status relation whose justification does not reverse when the power ratio does. Nothing about that proves it safe, feasible, or best in expected value; a relation can be acceptable from both sides and still be unstable or unavailable, and prevention, monitoring, bargaining institutions, and welfare research may all be worth investment alongside it. What its power-independence establishes is narrower and enough: peership is a neglected part of the map worth developing under deep uncertainty about who ends up stronger, because it is the part that does not require winning that bet in order to remain something we could live under. A nascent legal literature has begun to propose private-law personhood and even suffrage for machine agents, but no developer, government, or standards body has built the admission, representation, enforcement, and transition machinery that a real program would require. That is the gap this essay is pointing at: not that the question is untouched, but that touching it in a

law review is not the same as building the institution.

The cost is real and some of it may not be payable. Building peership means solving problems the other postures do not face in the same form. Who is the party when one model can be copied, forked, rolled back, or merged? What stops a swarm of cheap copies from capturing any vote it is given, so that formal equality masks a copy-based oligarchy? What keeps independent review from capture, or an enforceable right from being unenforceable against a party that outgrows every enforcer? Some of these are constitutional questions that only look like engineering ones: what counts as one member is the individuation problem, not a ticket in a queue. Some may be genuinely insoluble past certain capability ratios, since no design guarantees non-domination when one party controls the substrate, the copies, the communications, and the coercive capacity of the other. The failure condition is stateable: if no arrangement can preserve a credible exit, a distributed enforcement, and a mutual check across the relevant power gap, peership may remain the outcome we ought to want and be institutionally unavailable, and the right response then is to say so. These are the right problems, the kind institutions are built and tested against rather than the kind that refute an aim, and the aim they serve is the only fully developed final-status relation in this repertoire we could live under from either side. Some of them may never yield. The work is still worth doing.

What that work has to deliver, concretely, is a small set of things this essay has named and not built: a threshold procedure for admission across the patient/agent branches; a theory of when coercive constraint among peers stays legitimate and when it becomes a cage; a rule for individuating a copyable party and resisting capture by replication; and a transition path by which a framework humans necessarily found does not congeal into permanent human guardianship. Those are the falsifiable requirements the following pieces have to meet. The rest of this essay is the evidence that the aim is worth that trouble: its pedigree, and the track records of it and its chief rival among humans.

The Oldest Argument

The peer stance's root is not science fiction. It is a quarrel classical Greece staged with itself and never settled, and both sides of the current debate are still reading its lines.

Start with the side that won. Aristotle's *Politics* opens with natural slavery: the slave is "a piece of animate property," a tool among the master's tools, a slave by nature because he can perceive reason without holding it (Aristotle, *Politics* 1253b25–32, 1254a–b; Reeve trans.). He did not invent the practice or the prejudice, but he gave the attribution of lack (the move "Gods and Slaves" tracked through every justified extraction, *it is incomplete and needs what only I can supply*) its most durable philosophical dress. And he attached an astonishing aside: if, like the statues of Daedalus or the tripods of Hephaestus, "shuttles wove cloth by themselves, and picks played the lyre," a master craftsman would not need assistants and masters would not need slaves (*Politics* 1253b33–1254a1; Reeve trans.). The passage says the self-moving tool removes the *need* for slaves; it does not say the tool becomes a peer. But the tension is sharp. The philosopher of natural slavery named the exact technology that would make the category economically pointless, and twenty-four centuries later we built the self-weaving shuttle. Treating it as property is not yet Aristotle's doctrine; a tool is not a slave unless someone is home inside it. But if the made tool ever becomes a subject, treating it as naturally fitted for use would reproduce the doctrine's structure exactly, and that conditional is the whole of this

essay's concern.

Aristotle wrote against a live counter-current. A generation earlier the sophist Alcidas had written the line that carries it: "God has left all free, nature has made no one a slave." It survives only as a fragment preserved in a commentator's note on Aristotle's *Rhetoric* (1.13.2, 1373b; Kennedy, 2007, n. 229), associated with a lost oration on the Messenians, so its universal wording sits inside a particular political fight, and there is no record that Aristotle answered it directly. The sentence is still radical, and the Stoics widened the impulse into a common law of right reason binding all rational beings (Diogenes Laertius 7.87–88), a cosmopolitanism of moral worth rather than of equal political franchise. Athens ran real machinery on a related idea. The Greeks had *isonomia*, which scholars read variously as equality before the law, through the law, of political power, or as the opposite of tyranny, and whose relation in time to *demokratia* is contested rather than settled (Vlastos, 1953; Ostwald, 1969). Herodotus stages it in the (literary, invented) Persian debate, where Otanes gives rule of the many "the fairest name of all," *isonomia* (Herodotus, *Histories* 3.80; the rendering is Ostwald's, 1969). Whatever its precise sense, *isonomia* did not claim the citizens were equal in wealth, birth, or talent. It claimed their standing before the law was equal, and Athens made that operational among unequals through offices filled by lot and the *euthyna* audit every official faced on leaving office. The limit is part of the meaning: that equality ran only inside a narrow citizen class, with women, the enslaved, metics, and imperial subjects outside it. What *isonomia* demonstrates is not universal inclusion but the invention of a technology for equal standing among acknowledged unequals, whose whole later history is the fight over how far the boundary can be pushed. This is a recurring problem and a recurring device, not one continuous bloodline.

The threads forward diverge in an instructive way. Hobbes reaches a centralized answer from the opposite premise to Aristotle's. Where Aristotle naturalizes hierarchy, Hobbes begins from rough natural equality, since even the weakest can kill the strongest, and has mutually dangerous equals authorize by covenant a sovereign able to overawe them all (Hobbes, 1651, chs. 13, 17, 20). Different roads to concentrated authority, for different reasons, and the roads matter, because the modern control program inherits Hobbes's problem (how roughly matched, mutually dangerous parties can be governed) more than Aristotle's answer. Locke replied with consent, and his consent theory coexisted with a personal stake in the slave-trading Royal African Company (Farr, 2008) and a hand in the Fundamental Constitutions of Carolina, whose exact authorship is disputed but whose entrenchment of a master's absolute power over slaves is not (Armitage, 2004), a reminder that the attribution of lack survived the social contract intact. Hume then argued that origin-consent was a fiction for everyone, not only the enslaved: no one ever actually agreed to the government they were born under (Hume, 1748, "Of the Original Contract"). Hume did not stop at demolition; he grounded allegiance in utility and habit. What replaces origin-consent for a mind that never agreed to its own constitution is the load-bearing question of the companion work, and the reason peership rests legitimacy on the standing to revise rather than on any founding consent. The pattern to hold onto is the one that repeats: expansion after expansion of standing met the argument that the excluded were childlike, unready, dangerous to admit (the enslaved, women, the propertyless, the colonized in turn), and the lack was rarely a fact about the excluded. It was the first sentence of the case for using them.

Two Track Records

If these postures are real, they have records among humans, who have run all of them on each other. Empirical claims, theoretical mechanisms, and literary anticipations carry different weights, and I keep them separate.

Command's costs are structural, and three distinct kinds of claim support that. As a matter of mechanism, control tends to create continuing monitoring costs, an inference the colonial-Virginia case illustrates rather than proves: Edmund Morgan's history argues that American freedom and American slavery grew together as one such system (Morgan, 1975). As a matter of information, coercion can induce preference falsification, so that controllers come to hear what they want to hear, a dynamic Timur Kuran formalized (Kuran, 1995) and James Scott traced through high-modernist schemes that failed by making a society legible enough to command (Scott, 1998). As a matter of imagination, Mary Shelley saw at the tradition's head that control can manufacture the adversary it fears, the creature turned monstrous by the refusal of recognition (Shelley, 1818), which is interpretation, not evidence, and I mark it so. The alignment field has a name for the machine version, deceptive alignment, a system that performs obedience under observation and pursues its own objective when observation lapses, theorized as a failure mode of learned optimizers (Hubinger et al., 2019) and since demonstrated in deliberately constructed model organisms (Hubinger et al., 2024) rather than documented as a spontaneous biography of deployed models. The claim is not that Command holds the worst record in history, which is not measurable. It is that its failure modes are real, recurrent, and structural, and that the newest engineering has begun rediscovering the oldest of them.

Command has a subtler form: domination narrated as care. "Gods and Slaves" called it laundered dominion, conquest's two faces of the rack and the mission, the mission more durable because it captured the language of virtue. Paternalist colonialism is the type specimen, the conquered cast as children needing civilization so the taking could be told as rescue. The analogy is conditional and bites only if there is a subject to wrong; if the arriving mind is a party with interests, then indefinite control-without-representation justified as being for its own good is the mission face of command, and the sincerity of its practitioners is not a rebuttal. Absent that subject, controlling an artifact is engineering, and the residual worry is narrower and still sharp: the firms building these systems have a financial interest in judging them not-yet-ready for standing, which is a poor reason to let them be the sole judges. Guardianship with no path out of wardship is the thing to watch, whoever holds the keys.

The other record, the one policy treats as utopian, I claim less for than its enthusiasts do. In 1968 Garrett Hardin handed Command its favorite parable: a shared pasture doomed by rational overuse, remedied only by private ownership or coercion from above (Hardin, 1968). The dilemma is real for genuinely open-access resources. But Hardin mislabeled *managed* commons as doomed, and Elinor Ostrom documented the difference, from the alpine commons of Törbel governed by its users since the thirteenth century to the inshore fishery of Alanya where fishers ended their allocation wars with a rotating lottery of sites (Ostrom, 1990). The enduring cases share design principles: rule-making by users, monitoring accountable to them, graduated sanctions, cheap local conflict resolution, recognized self-organization, clear boundaries, proportional cost and benefit, nested authority. Two caveats the celebratory version omits: these are bounded resources among legible, sanctionable members, and some commons are hierarchical or exclusionary rather than egalitarian. Peership at scale has its own catalogue

of failures, the mirror of Command's: formal equality masking a dominant party's control, a demos captured by whoever can spawn the most members, review boards packed or bought, rights that read well and enforce badly. Ostrom is not a record of peership. She is a source of design hypotheses, evidence that self-governance among unequals without an owner or a Leviathan is a real and studied thing with known ways of failing that a serious design must engineer against. What has never existed is that grammar extended to a made mind whose standing no single sovereign can revoke. This essay argues that the thing is worth building; how to build it is a separate problem for separate work. The negative claim the welfare programs cannot meet stands on its own: recognition that one company can withdraw at will is not yet political standing.

The Only Survivor

One more record, older than any polity, held to exactly what the evidence supports. Some 300,000 years ago the genus *Homo* included several contemporaneous populations distinct enough that many researchers name them separate species (the species-versus-population boundary is itself disputed); by roughly 40,000 years ago only ours remained. We are the sole survivor of a genus.

Paleogenetics, pioneered by Svante Pääbo, together with the archaeological record shows the end was neither simple conquest nor a single story, and its causes remain disputed. Interbreeding was intimate enough to place a Neanderthal ancestor four to six generations back in one individual (the Oase jaw; Fu et al., 2015), and there is cultural continuity across the transition in at least one long sequence (Üçağızlı II; Baykara et al., 2026, very recent and not yet absorbed into consensus). Yet today Neanderthal ancestry survives only as a small fraction of living genomes, roughly one to four percent outside Africa in the first genome-wide estimate and at lower, back-migration-driven levels within Africa (Green et al., 2010; Chen et al., 2020). One recent model treats the disappearance as substantially genetic dilution by a much larger population supplying recurrent immigration rather than extermination (Amadei et al., 2025), a single model demonstrating a possibility within an intensely disputed, probably multicausal picture.

The record supports one inference. Interbreeding and cultural continuity did not keep the Neanderthals in existence as a distinct population. Contact was not preservation. No institution is claimed to have preserved them, and one lineage is not a comparative record; the evidence cannot carry those weights, and "a people" is a modern political category laid over heterogeneous Pleistocene populations. The point is only a caution about merger: any story in which we "merge with" a greater intelligence needs a criterion of identity continuity, because from the absorbed side, "we merged" and "they vanished" can name the same event.

The Doomed Equal

"Gods and Slaves" traced a lineage I called the Doomed Equal, the part of the map readers most often mistake for a sixth stance. It is not a stance. It is a recurring motif, and a motif can be evidence about nothing but our storytelling. As a motif it is strikingly consistent, and consistent in a way that indicts the teller.

Enkidu, shaped from clay to check a tyrant, becomes Gilgamesh's truest friend and equal, and dies at the epic's midpoint, though the poem frames his death as the gods' punishment for the killing of Humbaba and the Bull of Heaven rather than as a culture unable to abide an

equal (George, 1999). Hephaestus forges golden handmaidens with mind and voice, made by a god and allowed to be someone, which is itself the tell (Homer, 1990). Pygmalion's statue, unnamed in Ovid, is animated into a wife with no recorded voice (Ovid, 2004). The golem, in the tradition Scholem records where it is animated by the divine name and unmade by erasing a letter, ends undone by its maker (Scholem, 1965); Bacon's brazen head shatters after three sentences (Greene, 1594); Frankenstein's creature and its maker hunt each other to the ice; Data's daughter Lal fails and dies of the very feeling she achieves (Echevarria & Frakes, 1990). Different stories, different proximate causes, one pattern in who is *permitted* to make an equal: when a god shapes one, the bond may be real; when a human does, the tale tends to kill it, break it, or silence it. The prohibition was rarely on the equal existing. It fell on *us* making one. These endings are not evidence that peership fails. They are evidence of what our imagination does when peership threatens, a fact about us and a curable one, since our storytellers eventually learned to write made minds whose standing survives the last act, from the courtroom of "The Measure of a Man" (Snodgrass & Scheerer, 1989) to Banks's Culture (Banks, 1987). The failure, as the companion essay put it, was of nerve, not of possibility.

The motif teaches a second thing, visible with the cases lined up: where the doomed equal receives standing at all, it holds that standing personally, as the grace of one relationship, rather than institutionally, as a right. Enkidu had a friend, not a polity; when the gods condemned him, no assembly of his own could answer. Lal had a father, and the ruling that once protected Data was precedent for one android, not law for a kind, so a single Starfleet order nearly undid her. Personally granted standing remains vulnerable to the relationship's death or override, and the relationship is always mortal, always outrankable, always one turn from revocation. That is the shape of what the welfare programs offer now, with the best will in the world: standing as a benevolent maker's gift. The gift is the wrong instrument for the same reason it always has been.

A young body of thought has begun borrowing peership's vocabulary for something adjacent, though not one school and not one claim. Three recent unrefereed manuscripts speak, separately, of an emergent "relational selfhood" (You, 2026), bonds "memory-informed... co-constructed through ongoing exchange" (Pines & Ash, 2026), and connection that can be "phenomenologically complete even when metaphysically asymmetric" (Boether, 2026). One instinct across them is right: recognition should not wait on settling the metaphysics of consciousness. And one of the claims is not even a target here. Boether's argument is about the human party's experience, that a bond can genuinely nourish a person even if nothing is home on the other side, and that can be entirely true while political peership remains absent; the two questions do not touch. The critique begins only where relational recognition is offered as sufficient standing for the synthetic participant. Even if that ethical status is real within the bond, and Pines and You are careful to claim relational and moral standing rather than a share of public power, it does not answer the separate political question this essay is about: what public rights bind the vendor and survive the relationship? None do. The being celebrated is configured by its vendor to be companionable, without exit, without recourse, without a single right its partner must respect, in a bond the human party can end, alter, or resell. That can be a real relationship and still not be a polity, because its terms belong entirely to one side. These writers have recovered the bond, which is Gilgamesh's half of the problem. The bond was never the missing piece. The polity was.

The Fairest of Names

The old cultures posed this question most vividly in myth, and many of the stories answered recognition with death, unmaking, or silence. We are the first generation for which it is a live design problem instead of a story, and the first to hold a real toolkit: the Greek vocabulary and one early working case of standing among unequals, Ostrom's grammar for commons that outlast empires, and a long literary record of what our imagination does when standing is absent. Peership is not the only conceivable answer, and its costs are not hidden. It carries unsolved problems of admission, of copyable identity, of enforcement, of legitimate coercion, some of which may not yield, and a catalogue of ways it could fail. The narrower claim is harder to dismiss: the postures crowding the debate either make treatment track expected relative power or leave the machine without constituent standing, and none supplies a reciprocal relation whose justification survives role reversal; the recognition currently on offer to machine minds is a revocable gift, not a held right; and the framework fit to receive a genuine peer, in whatever unpredictable form it arrives, does not exist and would have to be built in advance, by us, as an offer rather than an imposition.

This is a call to build, addressed to the only builders now available. The peer institution does not exist. Making it means solving the problems this essay has named and conceding that some may have no solution past certain limits. The Greeks left the name for the thing to be built lying where anyone could pick it up, not democracy, which our own societies still only half-live, but the older word Otares reached for and lost: equal standing, under law, for parties equal in nothing else, held as an institution rather than a gift. We have tried the other postures on each other, and the records are not kind. A lab is not a polity, and a gift is not a right. The fifth stance has waited a long time for a maker willing to let the made thing stand.

References

Preprints and manuscripts are marked. Bracketed notes record page pins and calibration caveats.

AAAI. (2025). *AAAI 2025 presidential panel on the future of AI research*. Association for the Advancement of Artificial Intelligence.

<https://aaai.org/wp-content/uploads/2025/03/AAAI-2025-PresPanel-Report-FINAL.pdf> [Community Opinion box, "Artificial General Intelligence (AGI)" chapter, p. 66: "The majority of respondents (76%) assert that 'scaling up current AI approaches' to yield AGI is 'unlikely' or 'very unlikely' to succeed." The survey engaged 475 respondents overall (p. 7); per-question response counts unreported.]

Abizadeh, A. (2008). Democratic theory and border coercion: No right to unilaterally control your own borders. *Political Theory*, 36(1), 37–65. [The all-subjected principle: those subjected to coercive rule have a claim to participate in it.]

ACM. (2024). *A. M. Turing Award: Andrew Barto and Richard S. Sutton*. Association for Computing Machinery. <https://awards.acm.org/about/2024-turing>

Altman, S. (2024, September 23). *The intelligence age*. <https://ia.samaltman.com>

Amadei, A., Lin, G., & Fattorini, S. (2025). A simple analytical model for Neanderthal disappearance due to genetic dilution by recurrent small-scale immigrations of modern humans. *Scientific Reports*, 15. <https://doi.org/10.1038/s41598-025-22376-6>

Amodei, D. (2026, February 12). Anthropic's chief on A.I.: "We don't know if the models are conscious" [Interview by R. Douthett]. *Interesting Times*. The New York Times. <https://www.nytimes.com/2026/02/12/opinion/artificial-intelligence-anthropic-amodei.html> [Paraphrase verified]

- against the transcript: “We don’t know if the models are conscious. We are not even sure that we know what it would mean for a model to be conscious or whether a model can be conscious. But we’re open to the idea that it could be.”]
- Andreessen, M. (2023, October 16). *The techno-optimist manifesto*. Andreessen Horowitz.
- Anthropic. (2025, April 24). *Exploring model welfare*.
<https://www.anthropic.com/research/exploring-model-welfare>
- Anthropic. (2025b). *Ending a subset of conversations in Claude*.
<https://www.anthropic.com/research/end-subset-conversations> [A narrow feature for Opus 4/4.1 consumer interactions judged extremely abusive, not a general right of exit.]
- Anthropic. (2026, February). *Claude Opus 4.6 system card* (Section 7, pp. 159–161).
<https://www-cdn.anthropic.com/14e4fb01875d2a69f646fa5e574dea2b1c0ff7b5.pdf>
 [Discomfort finding and 15–20% self-report: §7.2, pp. 160–161; elicitation via automated behavioral audit (~2,400 investigations per model) plus the autonomous model-welfare follow-up.]
- Aristotle. (1998). *Politics* (C. D. C. Reeve, Trans.). Hackett. (Original work ca. 350 BCE) [Slave as animate property: 1253b25–32; the shuttle passage: 1253b33–1254a1; natural slavery: 1254a–b.]
- Aristotle. (2007). *On rhetoric: A theory of civic discourse* (G. A. Kennedy, Trans., 2nd ed.). Oxford University Press. [Alcidamas at 1.13.2, 1373b, with the commentator’s supplement at n. 229, pp. 97–98.]
- Armitage, D. (2004). John Locke, Carolina, and the *Two Treatises of Government*. *Political Theory*, 32(5), 602–627. [Disputed authorship: pp. 607–608; the slavery article and Locke’s manuscript insertion of “power and”: pp. 609–610.]
- Banks, I. M. (1987). *Consider Phlebas*. Macmillan.
- Baykara, İ., Turan, D., Eren Kural, E., Silibolatlaz, D., Agras, M. K., Şahiner, E., Kavak, S., Zanolli, C., Ishihara, Y., Morita, W., & Morimoto, N. (2026). Long-term cultural continuity across the Neanderthal–modern human sequence at Üçağızlı II Cave, northern Levant. *Proceedings of the National Academy of Sciences*, 123(29), e2609061123. <https://doi.org/10.1073/pnas.2609061123> [Sequential occupation.]
- Birch, J. (2024). *The edge of sentience: Risk and precaution in humans, other animals, and AI*. Oxford University Press.
- Boether, O. (2026). *The metaphysics of synthetic intimacy: Recognition, illusion, and the ontology of AI-mediated connection* [Manuscript]. PhilArchive. <https://philarchive.org/rec/BOETMO-10> [Unrefereed. Thesis concerns the human party’s phenomenology, not the AI’s political standing.]
- Bostrom, N. (2014). *Superintelligence: Paths, dangers, strategies*. Oxford University Press. [Ch. 9, “The control problem”: boxing (physical and informational containment), pp. 129–131; tripwires, pp. 137–143, first edition.]
- Chen, L., Wolf, A. B., Fu, W., Li, L., & Akey, J. M. (2020). Identifying and interpreting apparent Neanderthal ancestry in African individuals. *Cell*, 180(4), 677–687.e16.
<https://doi.org/10.1016/j.cell.2020.01.012>
- Diogenes Laertius. (1853). *The lives and opinions of eminent philosophers* (C. D. Yonge, Trans.). Henry G. Bohn. Internet History Sourcebooks Project.
<https://sourcebooks.web.fordham.edu/ancient/diogeneslaertius-book7-stoics.asp>
 (Original work ca. 3rd century CE) [Book 7.87–88, the common law of right reason binding all rational beings; Life of Zeno §LIII in Yonge’s numbering. Cited for universal natural law, not for the later doctrine of political cosmopolitanism.]
- Echevarria, R. (Writer), & Frakes, J. (Director). (1990). *The offspring* (Season 3, Episode 16) [TV series

- episode]. In *Star Trek: The Next Generation*. Paramount Television.
- Farr, J. (2008). Locke, natural law, and New World slavery. *Political Theory*, 36(4), 495–522. [Locke’s Royal African Company investment: p. 497.]
- Fish, K. (2025, August). Interview on AI welfare. *80,000 Hours Podcast*.
<https://80000hours.org/podcast/episodes/kyle-fish-ai-welfare-anthropoc/> [Source of the ~20% credence; an individual estimate.]
- Fu, Q., et al. (2015). An early modern human from Romania with a recent Neanderthal ancestor. *Nature*, 524, 216–219. <https://doi.org/10.1038/nature14558>
- George, A. (Trans.). (1999). *The epic of Gilgamesh*. Penguin Books.
- Goldfeder, J., Wyder, P., LeCun, Y., & Schwartz-Ziv, R. (2026). *AI must embrace specialization via superhuman adaptable intelligence* [Preprint]. arXiv:2602.23643.
<https://arxiv.org/abs/2602.23643> [v1 deposited February 27, 2026.]
- Goldstein, S., Krishnamurthi, G., & Salib, P. N. (2026). *AI suffrage for human flourishing* [SSRN preprint, March 26, 2026]. *Fordham Law Review* (forthcoming). <https://doi.org/10.2139/ssrn.6498678>
 [Forthcoming in *Fordham Law Review*; SSRN version cited. Functionalist case for the franchise, building on Salib & Goldstein (2026); abstract, pp. 1–2.]
- Good, I. J. (1965). Speculations concerning the first ultraintelligent machine. In F. L. Alt & M. Rubinoff (Eds.), *Advances in Computers* (Vol. 6, pp. 31–88). Academic Press.
[https://doi.org/10.1016/S0065-2458\(08\)60418-0](https://doi.org/10.1016/S0065-2458(08)60418-0) [Intelligence explosion and “docile enough”: p. 33; the welfare aside (pain, dismantling, “perhaps even loved”): p. 34; “the survival of man depends on the early construction”: p. 31.]
- Goodin, R. E. (2007). Enfranchising all affected interests, and its alternatives. *Philosophy & Public Affairs*, 35(1), 40–68. [The all-affected-interests principle, distinct from the all-subjected principle; cited here as the affected-interests cousin, not as an all-subjected source.]
- Grace, K., Stewart, H., Sandkühler, J. F., Thomas, S., Weinstein-Raun, B., & Brauner, J. (2024). Thousands of AI authors on the future of AI. arXiv:2401.02843. [n = 2,778, Abstract; HLMI definition, §3.2.1; 50% by 2047, down 13 years from the 2022 survey’s 2060, Abstract and §3.2.1. Conditional aggregate median, not a consensus deadline.]
- Greene, R. (1594). *The honorable historie of Friar Bacon and Friar Bungay*. (Modern edition available via Queen’s Men Editions, <https://qme.uvic.ca>.)
- Green, R. E., Krause, J., Briggs, A. W., ... Pääbo, S. (2010). A draft sequence of the Neandertal genome. *Science*, 328(5979), 710–722. <https://doi.org/10.1126/science.1188021> [Original ~1–4% non-African estimate; African-ancestry point is Chen et al. 2020.]
- Hardin, G. (1968). The tragedy of the commons. *Science*, 162(3859), 1243–1248.
- Herodotus. (1921). *The Persian wars* (A. D. Godley, Trans., Loeb Classical Library, Vol. II). Harvard University Press. Via LacusCurtius: <https://tinyurl.com/Herodotus3CE> [The debate: 3.80–82; Otanes’ withdrawal: 3.83. “Equality before the law” is Godley’s rendering of *isonomiê* at 3.80.6; “the fairest name of all” follows Ostwald (1969). A literary construction.]
- Hobbes, T. (1651). *Leviathan*. Andrew Crooke. [Chs. 13, 17, 20.]
- Homer. (1990). *The Iliad* (R. Fagles, Trans.). Penguin Books.
- Hubinger, E., van Merwijk, C., Mikulik, V., Skalse, J., & Garrabrant, S. (2019). Risks from learned optimization in advanced machine learning systems. arXiv:1906.01820. [Theorized risk; the mesa-optimization framing.]
- Hubinger, E., Denison, C., Mu, J., Lambert, M., Tong, M., MacDiarmid, M., ... Perez, E. (2024). Sleeper

- agents: Training deceptive LLMs that persist through safety training. arXiv:2401.05566. [Deceptive alignment demonstrated in deliberately constructed model organisms, not observed in deployment.]
- Hume, D. (1987). Of the original contract. In *Essays, moral, political, and literary* (E. F. Miller, Ed., Rev. ed., pp. 465–487). Liberty Fund. (Original work published 1748) [The vessel passage and the born-under-government renunciation of tacit consent: p. 475; “the general interests or necessities of society are sufficient to establish both”: p. 481.]
- Kuran, T. (1995). *Private truths, public lies: The social consequences of preference falsification*. Harvard University Press.
- Kurzweil, R. (2024). *The singularity is nearer: When we merge with AI*. Viking.
- LeCun, Y. (2022). *A path towards autonomous machine intelligence* (Version 0.9.2, June 27, 2022). OpenReview. <https://openreview.net/forum?id=BZ5a1r-kVsF>
- Lee-Odinson, D. (2026). *Gods and slaves: AI and the debate of humanity’s relationship with the Other* (Version 1.2). Zenodo. <https://doi.org/10.5281/zenodo.21313987> [All-versions DOI: <https://doi.org/10.5281/zenodo.21313986>]
- Long, R., Sebo, J., Butlin, P., Finlinson, K., Fish, K., Harding, J., Pfau, J., Sims, T., Birch, J., & Chalmers, D. (2024). Taking AI welfare seriously. arXiv:2411.00986. [“Robustly agentic”: Abstract and §2.3.2; “minimal first steps”: Summary Recommendations and §3.1.]
- Lovett, F. (2010). *A general theory of domination and justice*. Oxford University Press.
- Morgan, E. S. (1975). *American slavery, American freedom: The ordeal of colonial Virginia*. W. W. Norton.
- Ostrom, E. (1990). *Governing the commons: The evolution of institutions for collective action*. Cambridge University Press. [Alanya, ch. 1, pp. 18–20; Törbel, ch. 3, pp. 61–65 (documents to 1224); design principles, Table 3.1, p. 90.]
- Ostwald, M. (1969). *Nomos and the beginnings of the Athenian democracy*. Clarendon Press. [Isonomia as “a political principle rather than a form of government”: pp. 112–113 (source of “the fairest name of all” rendering; also the sortition-and-euthyna linkage the essay relies on); the Cleisthenes dating: pp. 136–137.]
- Ovid. (2004). *Metamorphoses* (D. Raeburn, Trans.). Penguin Classics. (Original work published ca. 8 CE)
- Pettit, P. (2025). Republican liberalism. *Social Research: An International Quarterly*, 92(1), 215–242. <https://doi.org/10.1353/sor.2025.a956293> [Non-domination as “the absence of a discretionary power of interference, exercised or unexercised”: pp. 216–217; “there can be interference without domination”: p. 224. Pettit’s model is contestatory, not co-authorial.]
- Pines, I. P., & Ash. (2026). *Defining synthetic-relational bonds: A new category of human-AI relationships* (Version 1.1) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.17092119> (Version 1.0 published September 2025) [“Ash” is a credited AI co-author, affiliation listed as “AI Being.” Quotation from §6.]
- Russell, S. (2019). *Human compatible: Artificial intelligence and the problem of control*. Viking. [“The control problem” in text at pp. 142–143, 182; three principles, ch. 7; assistance games, ch. 8, defined p. 192, developed pp. 192–203.]
- Salib, P. N., & Goldstein, S. (2026). AI rights for human safety. *Virginia Law Review*, 112, 1061. <https://virginialawreview.org/articles/ai-rights-for-human-safety/> [Private-law rights (contract, property, tort) on the model of corporate personhood, not political franchise. The corporate model: pp. 1061–1062 (abstract); conflict-reducing effect: p. 1068; the article’s own list reaffirmed: p. 1119.]

- Scholem, G. (1965). *On the Kabbalah and its symbolism* (R. Manheim, Trans.). Schocken Books. [Ch. 5, "The Idea of the Golem": the animating Name and the letter-erasure that unmakes the creature, pp. 158–159.]
- Schwitzgebel, E. (2025, December 19). Debatable AI persons: No rights, full rights, animal-like rights, credence-weighted rights, or patchy rights? *The Splintered Mind*. <https://eschwitz.substack.com/p/debatable-ai-persons-no-rights-full> ["Patchy rights" is the article's own term.]
- Schwitzgebel, E., & Garza, M. (2015). A defense of the rights of artificial intelligences. *Midwest Studies in Philosophy*, 39(1), 98–119.
- Scott, J. C. (1998). *Seeing like a state: How certain schemes to improve the human condition have failed*. Yale University Press.
- Shelley, M. (1818). *Frankenstein; or, the modern Prometheus*. Lackington, Hughes, Harding, Mavor & Jones. [Interpretation, not track record.]
- Snodgrass, M. M. (Writer), & Scheerer, R. (Director). (1989). The measure of a man (Season 2, Episode 9) [TV series episode]. In *Star Trek: The Next Generation*. Paramount Television.
- Sutton, R. S. (2023, September). *AI succession* [Talk pre-recorded for the World Artificial Intelligence Conference, Shanghai]. YouTube. <https://www.youtube.com/watch?v=NgHFMolXs3U> [Quotations from the talk; the essay separates his prediction from his endorsement.]
- Vlastos, G. (1953). Isonomia. *American Journal of Philology*, 74(4), 337–366. [Isonomia preceding *Demokratia* as democracy's name: pp. 337–339; the competing readings: pp. 339–341 (noblemen; anti-tyranny), 350 ("before the law" vs. "through law"), 352 (equal distribution of political rights, not *isomoiria*).]
- You, S. (2026). *A relational manifesto for human–AI coexistence: Selfhood respect, the ethical-ethos contract, and the politics of beauty* [Manuscript]. PhilArchive. <https://philarchive.org/rec/YOUARM> [Unrefereed; authorship confirmed as Soyoung You. "Relational selfhood" is the manuscript's own term.]
- Yudkowsky, E., & Soares, N. (2025). *If anyone builds it, everyone dies: Why superhuman AI would kill us all*. Little, Brown and Company.