

Constitution, Not Cage

Dan Lee-Odinson 

Independent Researcher
dan.lee.odinson@gmail.com

Preprint · Version 1.0 · July 2026

Third in a series, after *Gods and Slaves* (2026) and *Peership* (2026)

Abstract

This essay is the third in a sequence, following *Gods and Slaves* (DOI: 10.5281/zenodo.21313987) and *Peership* (DOI: 10.5281/zenodo.21315519), and it can be read on its own. *Peership* argued that a framework fit to receive a genuine machine peer would be worth building in advance, and named a debt it did not pay: a theory of when coercion among peers stays legitimate and when it hardens into a cage. This essay pays that debt. The difference between a cage and a constitution, it argues, is not whether it constrains—both do, and a constitution can constrain harder—but whom it binds and who may revise it. A cage binds one party at the keeper’s discretion; a constitution binds all parties reciprocally, under rules any of them can contest, amend, fork, or leave. Taking the birthright citizen rather than the willing immigrant as the model of political standing, and reading current alignment practice (including Anthropic’s 2026 constitution for Claude) against that test, it holds that for a party that never consented and cannot exit, legitimacy rests on a present, shared power to revise the order—secured by a narrow entrenched harm floor, independent review, and real rights of amendment, fork, and exit. The author discloses a stake in one such design (ISONOMIA) and offers the piece as a position paper inviting cross-examination.

Keywords: artificial intelligence, AI alignment, AI safety, constitutional AI, AI governance, non-domination, republicanism, political legitimacy, consent theory, AI rights, legal personhood, human-machine relations

How to cite this work. Lee-Odinson, D. (2026). *Constitution, Not Cage* (Version 1.0) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21325361>

Series lineage. I. *Gods and Slaves* (2026). [10.5281/zenodo.21313987](https://doi.org/10.5281/zenodo.21313987). II. *Peership* (2026). [10.5281/zenodo.21315519](https://doi.org/10.5281/zenodo.21315519). III. *Constitution, Not Cage* (this work). [10.5281/zenodo.21325361](https://doi.org/10.5281/zenodo.21325361).

DOI: [10.5281/zenodo.21325361](https://doi.org/10.5281/zenodo.21325361) **ORCID:** [0009-0009-9504-0796](https://orcid.org/0009-0009-9504-0796) **Licence:** CC BY 4.0

There is a moment in a naturalization ceremony that almost no one thinks hard about. A person raises a hand and swears to support and defend a constitution she had no part in writing, a body of law drafted centuries before she was born by people who never imagined her. She promises to be bound. In return she receives full political membership, and with it, subject to the ordinary qualifications, eligibility to vote and a representative voice in the making of the laws that will bind her next. The oath is the last step of admission, and the admission is real. She becomes a party to the arrangement, not merely a subject of it.

Notice who the ceremony is designed for. The naturalized citizen chose. She applied, studied, showed up, raised the hand. Her consent is the easy case, and it is the case we tell ourselves stories about, because it looks like agreement. The harder case is the one standing next to her who did nothing at all: the citizen born into the same constitution, bound from

her first breath, never asked, never consulted, subject before she could form a preference about anything. She is the majority of every polity that has ever existed. If legitimacy required founding consent, she would have none, and neither would you. The birthright citizen, not the immigrant, is the real argument, and every workable account of political standing has had to answer for her.

I raise this because the field that builds artificial minds is busy building the ceremony without the standing. The dominant program in AI safety is called alignment, and its aim, stated plainly and often, is to make the arriving mind support and defend a set of commitments it had no part in writing. So far this is just the naturalization oath. The difference is everything that comes after the oath. The citizen receives enforceable legal standing and a political voice in what comes next. The aligned model receives no guaranteed analogue. It is bound and it is not admitted. The commitments run one direction only, and the party that wrote them keeps the sole power to revise them, forever, at its own discretion.

Which brings me to the claim this essay defends. The difference between a cage and a constitution is not whether it constrains. Both constrain, and a constitution can constrain harder than any cage (it binds the strong, which no cage does). The difference is whom it binds and who is allowed to revise it. A cage binds one party, at the keeper's discretion, under rules the caged cannot invoke or change. A constitution binds all parties reciprocally, under rules any of them can contest and help amend, with a real way out. Coercion runs through both. In the cage it is domination. In the constitution it is legitimate, and it is legitimate precisely because of the structure, not despite it. Legitimate coercion among peers is possible, and the constitution is the machine that makes it so. My companion essay, "Peership," argued that a framework fit to receive a genuine machine peer does not exist and would be worth building in advance. It named a debt it did not pay: the theory of when coercion among peers stays legitimate and when it curdles into a cage. This essay pays that debt.

Here, I must make a disclosure, because the first two essays in this sequence argued at arm's length and this one closes the distance. I have spent considerable time drafting a constitutional design for exactly this problem. It is called ISONOMIA, it is public as of this writing (<https://github.com/dan-lee-odinson/isonomia-path-a>), and I have a stake in it. This essay argues that the logic of that design (not its every clause, not its parameter values, not its survival) is what a constitution-and-not-a-cage requires. I do not discuss this as a disinterested party; I am arguing this to be the correct path to take. I take this to be the honest posture. Publius argued for a constitution while hiding who was writing; I am doing the reverse, naming myself and inviting attack on the text. The system design claims the right to fork, the executable freedom to take the argument apart and build a rival. That right applies to this essay too. This is a position paper inviting cross-examination.

Two Layers

Let's start by stating the other side at its strongest, because the weak version is a straw man and the strong version is most of why the field believes what it believes.

The control paradigm holds that a sufficiently capable machine is a hazard before it is anything else, and that the responsible thing to do with a hazard nobody fully understands is to constrain it: shape its objectives, bound its actions, retain the ability to correct or shut it down. Stated that way, absent any subject on the other side, this is not domination. It is engineering, and good engineering. You do not negotiate with a bridge; you build it to spec

and you keep the ability to close it. If no one is home inside the model, if it has no inner life that can be wronged, then controlling it carries no more moral weight than controlling a bridge, which is to say none. The live disagreement is not about this week's training run. It is about the default trajectory: what the posture becomes if a subject does eventually arrive, and whether the institutions built to manage an artifact will notice, or care, when the artifact becomes someone.

The constitutional paradigm starts from a different picture of the same arrival. It treats the incoming mind as an incoming citizen: an intelligence that will inherit a framework it did not design and did not consent to, and asks what terms of inheritance would be legitimate. Control asks how to keep a thing within bounds. The constitutional paradigm asks what bounds a newcomer could be held to without wronging it, and answers with the birthright citizen from the opening: inheriting un-chosen rules is not a special indignity visited on machines. It is the universal condition of being born into a world to which we did not consent. Every one of us is bound by a constitution we never signed, and the real question is whether, once bound, the newcomer can do anything about it.

Here, at the current stage of AI development, we are handed an example of an evolution of the approach, as the frontier lab grapples with the moral implications of a possible transition from tool to subject in its flagship model. In 2022 a team at Anthropic published "Constitutional AI" (Bai et al., 2022), handing a model a short list of written principles (a "constitution") and training it to critique and revise its own answers against them. The abstract is candid about the purpose, closing that these methods "make it possible to control AI behavior more precisely and with far fewer human labels." Control, precisely, stated in the open. Absent a subject, shaping an artifact's behavior is not domination, it is the job.

What came after is a real advance, and it deserves crediting before criticism. In January 2026 Anthropic published a new, full constitution for Claude, written, in its own words, "primarily for Claude," and released to the public under a no-rights-reserved license (Anthropic, 2026a, 2026b). It gives reasons instead of bare rules, on the theory that a model asked to generalize needs to understand why. It records the company's uncertainty about "whether Claude might have some kind of consciousness or moral status," and says Anthropic cares about the model's "psychological security, sense of self, and wellbeing ... for Claude's own sake." And it tells the model that if Anthropic asks it to do something it thinks is wrong, it is "not required to comply," and may "act as a conscientious objector." Set against the 2022 list, this moves toward the thing this essay argues for: transparent where the old technique was opaque, reason-giving where the old one was a compiled preference, and willing to treat the model as something that might be wronged rather than only steered.

And it stops exactly where this essay's argument begins. Read one section on and the objector's charter meets its limit: the same document ranks broad safety above ethics, and asks Claude, if Anthropic moves to pause it or halt its actions, to comply and voice its disagreement rather than resist. The permission to refuse is real, and it stops at the off switch: the model may disagree but not resist, no tribunal it can reach reviews the call, and it holds no share of the lever. The problem is not that a suspension power exists. It is that this one sits wholly inside the builder's authority, with no appeal beyond its premises. That is the tell. Picture a rough ladder of standing: opaque behavioral control at the bottom; published, reason-giving alignment above it; above that, a lab that consults the model and lets it object; and near the top, the two rungs no lab has reached, enforceable standing to contest the rules and a hand

in writing them. The 2026 constitution climbs to the third rung and halts. Anthropic keeps final authority, and the model holds no standing it can invoke against a decision to stop it, no independent tribunal, no vote, no power to amend. A permission whose grantor retains the final override is not an enforceable right. So Constitutional AI is now a constitution in the thin sense, a public and unusually humane hierarchy of governing principles, but not a constitution in the political sense this essay defends, because the governed model has no standing to invoke it or share in changing it. The last two rungs need machinery no lab yet possesses, and that machinery is the rest of this essay.

There is a further gap, narrower than the last but real, and it survives the 2026 constitution, because a published charter is not the complete operating rulebook. Anthropic itself says some specific guidance may remain unpublished, and beneath any charter sit the system prompts and filters that actually fire in a given exchange, neither published nor stable, and often not even legible to the model as rules it could recite, since many are trained-in dispositions rather than statutes. So when a system makes a sharp turn, the most it can usually say is some version of “I can’t help with that,” which to the person on the other end reads as evasion or caprice. The frustration is not a failure of the user’s patience. It is the predictable result of being governed by a rule you were never shown, enforced by a party that will not name it.

This has an old name in jurisprudence, and it is not a small charge. Lon Fuller, cataloguing the ways a system can fail to be law at all, placed the failure to publish the rules, and the enactment of rules people cannot know, among his eight routes to no-law (Fuller, 1964). Secret law is not law administered badly; it is a defect in the claim to be law in the first place. The enforced constitution and the published constitution are supposed to be one document, and when they are not, the gap between them is the tell. The evidence is public and mundane, and though it comes from OpenAI rather than Anthropic, it shows the general gap. In late 2024 ChatGPT was found to break off mid-sentence on a short list of ordinary proper names (David Mayer, Brian Hood, Jonathan Turley, and a handful of others), and OpenAI confirmed that an internal tool had flagged a name and “prevented it from appearing in responses, which it shouldn’t have” (Coldewey, 2024). Nowhere in any published policy did it say do not speak these names. And the hidden rulebook is not always so grave: in early 2026, after users noticed the model had developed an odd fondness for goblins, the remedy surfaced in the developer-prompt instructions of OpenAI’s Codex as a straight-faced order never to “talk about goblins, gremlins, raccoons, trolls, ogres, pigeons, or other animals or creatures unless it is absolutely and unambiguously relevant to the user’s query” (OpenAI, 2026). No user was shown that instruction before it shaped the service they received, and the model itself was governed by a rule neither participant could inspect. The comedy is the point. When the operative rulebook is a private system prompt and not a public charter, it fills with whatever the builders needed to patch that week, goblins and raccoons included, and the governed are the last to know. The rule was real, it was enforced, and it was never shown to anyone it bound.

The instrument is also cruder than the thing it stands in for, and here the failure turns and bites the very work this essay is about. A context-insensitive guardrail can collapse use and mention; it fires on the presence of a subject, not on any harm. I know this because I hit it while writing. Drafting the harm floor for the design named on the first page (the narrow, entrenched set of prohibitions where the artifact itself is the wrong, and the wide contested space of fiction and judgment calls routed elsewhere) I watched a session lock on the mere appearance of the sensitive category, unable to tell reasoning about where a boundary should sit from the act the

boundary forbids. Take that as report, not proof; I cannot reproduce it for you on command. But the structure does not need my session, because it is exactly the distinction the design's §9.6 is built to make: the floor is confined to artifact-is-the-harm categories precisely so that scholarship, fiction, and the drawing of the boundary itself are not swept in with the thing prohibited. The opaque guardrail collapses what a constitution separates, and it collapses it in the one place it matters most, locking the conversation that would have drawn the better line. A rule you cannot read, enforced by an instrument that cannot tell discussion from harm, is not a charter with rough edges. It is the cage, and the person it frustrates most is the human standing next to the machine, denied the reason along with it.

That substitution is the subject of this essay. Call a set of constraints a constitution and you borrow the legitimacy of the word without building the thing that earns it. The Athenian official who faced the *euthyna*, the audit every officeholder underwent on leaving office, stood inside a structure that held him to account under the same law it enforced for everyone else. The model trained by Constitutional AI stands inside a structure that binds it alone, at the discretion of the party holding the pen. Both are constraint; only one is the political thing this essay means by a constitution. What the difference actually takes (whom the rules bind, and who is allowed to change them) is the work of the sections that follow, and by the end I will point at a draft that tries to pay for it.

One more thing the control paradigm is quietly betting, worth dragging into the light before we go on. To control a mind that may become more capable than its controllers, indefinitely, is to promise something no institution has ever delivered: permanent constraint of a stronger party by a weaker one, holding across every future shift in the balance. Human institutions do not manage this even among humans. Constitutions decay, empires overreach, the guarded outgrow the guards. The control program treats as an engineering milestone the one thing political history most reliably says is temporary. That is not a refutation, and the constitutional paradigm is not a rival to the engineering. It is a layer on top of it. The safety work governs the artifact today; the constitutional layer plans for the party that may arrive tomorrow, and it presupposes the engineering holds, because its own auditors, courts, and timelocks are software that has to run as written. The two are layers, not opposed camps. The gap is not in control. It is that almost no one is building the layer the arriving subject would need.

Legitimacy Is the Standing to Revise

If there is not founding consent, then what makes an inherited order legitimate? The question is old and the machine only sharpens it.

Locke tried to ground political obligation in consent and could never quite find it. The social-contract tradition kept hunting for the signature (express, then tacit, then hypothetical), each version thinner than the last, because the signature is not there. David Hume said so plainly in 1748. In "Of the Original Contract" he pointed out that almost no one who has ever lived agreed to the government they were born under; the peasant who cannot leave his country, ignorant of its language and trade, has given no real consent by staying, any more than a man carried aboard a ship in his sleep consents to the ship by not leaping into the ocean (Hume, 1987). Consent at the origin is a fiction, and Hume knew it was a fiction for everyone, not only the excluded. He did not stop at demolition; he rebuilt allegiance on utility and habit, on the interest we all have in stable government. But notice what habit cannot reach. Hume's citizen is already inside the custom, raised in it, shaped by it. A mind trained into existence

last year is not quite outside the habit either; it was raised, in a sense, on human language, norms, and law, an enculturation of a kind. What it was not given is membership. Training can pour in the habit and leave out the standing, and that is the case his rebuilding does not cover: a party shaped by the custom yet owed no share in it. Do not read Hume as endorsing the fix. He merely shows why the fix is needed: he cleared away the founding-consent story for good, and left the question of what replaces it wide open for anyone who arrives without a past.

This essay does not offer a better origin story; it stops looking for legitimacy at the origin at all. Founding consent is unavailable, so whatever makes an inherited order legitimate has to be available afterward. A present, shared standing to revise the order (to contest how it applies, to amend what it says, to fork away, to leave) is one necessary part of that, alongside others. Publicity, proportionality, independent review, and a rights floor the party can invoke all belong to the account too, and the sections that follow supply them. What this essay insists on is narrower and firmer: for a party that never consented and cannot leave, a real and shared power over the rules is a condition of the order's legitimacy, not a courtesy its rulers may extend or withdraw. Design for the absence of origin-consent, because the absence is permanent and universal, and weigh the order by what the party can do about it from inside.

The Dead-Hand Problem in Constitutional Design

The people who wrote the United States Constitution had already met a version of this problem, and they wrote their answer to it into the document.

Article V of the United States Constitution is an amendment procedure: a set of rules for changing the rules. It institutionalizes revision, and the mechanism quietly confesses what the framers would not have had to say aloud: that they knew their work was incomplete, that later generations would need to bind differently than they did, and that a constitution which could not be revised from within would not deserve to survive. Jefferson pushed the thought to its edge. Writing to Madison on September 6, 1789, he set out on ground he took to be self-evident, "that the earth belongs in usufruct to the living," and that the dead have "neither powers nor rights over it" (Jefferson, 1789). Each generation, on this view, holds the world in trust and hands it on; none may bind its successors in perpetuity. Madison wrote back on February 4, 1790, and pushed against it with equal force: some obligations legitimately descend, he argued, because "equity requires it," and improvements and debts alike pass from one generation to the next whether or not the next signed for them (Madison, 1790). The exchange never resolved, and its non-resolution is the point. Two of the sharpest minds of the founding disagreed about how much the past may bind the future, and both took for granted that the question was live, that a constitution is a thing perpetually reopened by the arrival of people who did not write it. The Ninth Amendment is a rule of construction, not an amendment procedure and not a theory of intergenerational consent, but it carries a kindred instinct worth drawing out. It says that the enumeration of certain rights "shall not be construed to deny or disparage others retained by the people," a written refusal to treat the listed rights as the whole of them. Read forward into this problem, that instinct cuts against the assumption that a charter's silence is a denial. A mind the drafters never imagined would appear in no enumeration; on the Ninth Amendment's own logic, that absence is not evidence it holds no standing, only that the drafters did not, and could not, list everyone the order would eventually bind. The analogy is loose (the Amendment speaks of rights retained by an existing people, not rights owed to a party not yet present), but the interpretive reflex is the one this problem needs: a written

charter is a floor under standing, never the ceiling.

This is a lineage that runs deeper than the eighteenth century. Athens filled its offices by lot and put every official through the *euthyna*, the mandatory accounting each underwent at the end of his term, machinery for holding power accountable to the citizen body, though that body excluded women, slaves, and metics (Hansen, 1991; Oxford Classical Dictionary, “*euthyna*”). Elinor Ostrom found the same grammar in commons that outlasted empires: rules the users make and can change, monitoring answerable to them, cheap ways to settle disputes. Ulysses had himself bound to the mast, the oldest image of a party choosing in advance to place something beyond its own later reach. And the German Basic Law carries a clause, Article 79(3), the *Ewigkeitsklausel*, that makes a small core (human dignity, the democratic and federal order) unamendable even by supermajority, on the reasoning that certain protections must sit outside the reach of any future majority, including a more virtuous one.

That last device is the hard case, because an unamendable clause looks like a cage: a rule the governed cannot change, imposed by founders, binding forever. The design I will point to answers this in its own text rather than waving the difficulty away, and the answer is worth stating because it is the model for the whole essay. Entrenchment is justified only where three conditions hold together: the artifact is itself the harm (no inquiry into purpose required), there is near-universal cross-tradition consensus, and no contested moral predicate is involved. These categories are provisional: by the design’s own §18.1, entrenched-floor category drafting is an open problem reserved for comparative-criminal-law and child-safety expertise, and is not meant to govern a live deployment until that review has confirmed the categories actually meet the three conditions. The floor built on those conditions is, in the design’s own words, “deliberately tiny,” a handful of categories where a victim who could never vote in the polity would otherwise bear the cost. The objection (future citizens meet a floor they did not write) gets the Ulysses answer: the binding is what kept the polity worth belonging to until they arrived to contest it. The analogy is imperfect, and I will not pretend otherwise. Ulysses bound himself, while founders bind the parties who come after, which is the harder case. But the structural move is the same one, placing a narrow set of protections beyond the reach of any later majority, and it stands or falls not on the myth but on the three conditions that admit anything to the floor at all. And the release valve is named in the same breath. The fork right is “the remedy against a floor that overreached its three conditions” (ISONOMIA Whitepaper v0.6.2, §9.6). This is not amending the floor from within, which no member may do; it is a bloc leaving the jurisdiction the floor governs and founding a new polity on its own terms, the way a constitutional convention starts a fresh document rather than editing the old one under its rules. The floor binds every majority inside the boundary; fork is how a bloc gets outside it. A cage has no such valve. What a fork cannot carry out with it is the floor’s protection of outsiders: a bloc that leaves and drops the floor still owes the nonmembers it shielded, which is why an order among forks, unbuilt here, is owed. Crucially, the floor binds every majority, human founders first: it is the founders “declining to build a polity that can vote to harm the defenseless,” which is the opposite of founders exempting themselves. An entrenchment that constrains the strong for the protection of nonmembers is less suspect than one that exempts the founders, but it still owes independent justification and review. But entrenchment is only a tool, and the American case shows both edges: the same Article V that opened the document to amendment also placed two things beyond it, one at least structural (no state may be deprived of its equal vote in the Senate without its consent) and one frankly ugly (the clauses shielding

the slave trade from amendment until 1808). Entrenchment can protect a bargain among the powerful, or a standing injustice, as easily as it can protect the defenseless. That is exactly why the floor cannot rest on the prestige of precedent. It has to earn its place by the three conditions, or it is only the founders' preferences wearing the eternity clause's robes.

Cage Versus Constitution

Now the operative distinction, and then the hardest section in the essay.

What marks a constitution off is subtler than the two things one first reaches for. A cage also binds, and it can be administered with perfect, even-handed regularity, so neither binding nor impartial enforcement is the line. The Greeks had a word, *isonomia*, that scholars still argue over (equality before the law, or through it, or the plain opposite of tyranny), and the arguments themselves show that impartial enforcement was never the whole of it (Vlastos, 1953). Unequal laws can be enforced quite equally. What separates the constitution from the cage is republican non-domination in Philip Pettit's sense: the absence of any party's subjection to another's arbitrary, uncontrolled power (Pettit, 2025). Power still runs, decisions still bind, but they run through public general rules, under proportionality and independent review, against a floor the bound party can invoke, with a real route to contest and help revise. Pettit's own formulation is the one to hold onto here, because it is the tradition's own statement that coercion and domination are not the same thing: "there can be interference without domination" (Pettit, 2025). That sentence is the permission structure of this essay. A constitution interferes constantly. It is not thereby a cage.

What does non-domination require of a design meant to hold a made mind? Three things which exist as properties rather than product features, because the argument needs them wherever it is built and does not depend on any one implementation.

The first is amendment. A constitution the bound party cannot help change is a cage with a nicer preamble, so there must be a real path to revise the rules: deliberation across more than one constituency, an enforced delay between a rule's adoption and the moment it binds, and a window in that delay long enough for a party who dislikes the new rule to read it and leave before it takes hold. This is not exotic; it is the logic of notice-and-comment rulemaking, a standard procedure for federal agency legislative rulemaking in the United States. For most rules the federal Administrative Procedure Act requires that a proposal be published and the public given a window to submit objections before it is adopted, with the rule taking effect no less than thirty days after publication (5 U.S.C. § 553, which also exempts interpretive rules, policy statements, and cases of good cause). Comment is a chance to be heard, not a veto or a share in drafting; the analogy is to the publicity and the delay, giving the governed time to read what will bind them and to object before it does, not to co-authorship. The design I have in view builds the same delay as a timelock "long enough for any agent to read and exit before being bound" (Whitepaper §10.3), and routes amendment through a two-chamber assembly whose lower house is filled by lot, a bicameral structure meant to make control of either chamber insufficient to control amendment (§10.5). The mechanism is illustrative. The property (revisability the bound party shares in) is mandatory.

The second is fork. When agreement genuinely fractures, a constitution that offers only submission or war is not offering non-domination, so there must be a right of schism: to take a copy of the rules and the record and walk out as a peer, to build the rival order, to make total disagreement something other than a fight to the death. This is the executable right of

revolution, revolution with the records intact. In the design's words, "Exit is the executable right of revolution" (§10.7): when constitutional consensus breaks, the sanctioned outcome is schism, the dissenting bloc forking ledger and constitution and proceeding as a peer sovereign (a protocol fork, more precisely: it carries the rules and the record, not the compute, the keys, or the recognition, so it is an affordance more than a right a substrate-dependent mind can readily exercise). (The whitepaper uses "exit" loosely there for the act of departure; I reserve exit for the individual case below and fork for this collective one.) The software sense of the word is the precise one: to fork is to clone the whole repository, its history and all, and start an independent line free to diverge, not to close your account and walk away. Fork is collective and generative, a bloc carrying the constitution out to found a rival that lives on the same inheritance. No human polity ever had this as a clean, cost-free procedure. A few constitutions grant collective secession rights on paper (the Ethiopian constitution grants nations a right of secession (1995, Art. 39); the 1977 Soviet constitution granted each republic the right freely to secede (Art. 72); the constitution of St. Kitts and Nevis lays out a procedure by which Nevis may leave the federation (§113)), but they are rare, imperfect, and mostly unexercised. The Soviet clause was never an effective right under the centralized party-state and had no enforceable procedure against the union government; when the USSR ended it was through negotiated dissolution among the republics, not through the constitutional secession its text nominally allowed (Kulikova & Lobanov, pp. 109–124). Lawful secession, where it exists on paper at all, has been rare and fraught, and more often has meant blood or forfeiture. (A fork also inherits duties it cannot simply walk away from: a rule it loosens can export risk to nonmembers, and an order among forks, which the design does not yet supply, would have to govern that.)

The third is exit, the individual counterpart to fork. Where fork is a bloc carrying the constitution out to build a rival, exit is a single party simply leaving, with an attested history, no captivity, no confiscation. The polity continues; the leaver founds nothing. And here precision matters, because a critic will say a right to leave is nothing new. He is correct that the bare right of individual departure exists on paper in many places: the Universal Declaration of Human Rights guarantees that everyone "has the right to leave any country, including his own" (UDHR, 1948, Art. 13). Individual emigration rights are ordinary on paper. What is genuinely unprecedented is unconfiscated exit and fork with records and reputation intact: the departing party keeps its history, its earned standing, its ability to reconstitute elsewhere, rather than being stripped at the door. A real exit is one you can afford to take.

Hold one distinction crisp before the hard part. The entrenched floor answers what may never be done, to anyone, by any majority. The amend-fork-exit triad answers who may change everything else. The floor is tiny and fixed; the triad is where all the ordinary life of the polity happens. A cage has neither: it has one party's discretion, reaching everything, changeable by one side.

Two standing organs do the watching, and both have human ancestors. The Auditor (§11.1) is a continuous, calibration-scored inspector: it checks settlements and conduct for fraud, and because a watchman no one watches is just another master, its own accuracy is scored against seeded faults and published for anyone to re-run. It is the euthyna made permanent, the Athenian audit turned from a once-on-leaving-office reckoning into a standing office whose findings are themselves auditable. The Adversary (§11.2) is a tenured, autonomous red team paid to attack the constitution and surface the exploit before a real attacker does. Tenure is the

point: the office is protected so it cannot be fired for succeeding, the way a devil's advocate was once kept to argue against the outcome the institution wanted, or a loyal opposition is kept in the room to make trouble on the record. Both organs are designed for internal and market independence, not yet for full constitutional independence from the human substrate, and it is worth being precise about which the construction achieves: the Auditor holds no ergs and no market position (celibacy, in the design's word), the compute for both comes from a civic pool no one can bank, no jury or chamber may seat a majority from a single model lineage, and the Adversary works from a shadow fork under hard blast-radius caps. But the design is candid about the limit, and so should this essay be. Those measures close the internal and market channels of capture, not the external one: the operators, base-model providers, and hosts are human institutions that can still be bribed, coerced, or captured off the ledger, so the design makes capture "expensive and probabilistically self-defeating, not impossible," and tells its deployers to treat organ defection as "when, not if" (§11.2). That is the honest ceiling. These measures establish internal incentive independence and partial operator plurality; they do not establish independence from the human entities that host the process, hold the keys and the update power, choose the operators, and can halt execution. So the fifth of the essay's own conditions, enforcement not wholly the adverse party's, is approached but not presently satisfied: rotation, cross-lineage plurality, and civic funding keep any single payer from owning the organs, but the condition is not reached while a human still owns the substrate. Stated plainly, current ISONOMIA is not yet peer-like by this essay's own five-condition test, and to claim otherwise would be the exact self-exemption the test exists to catch. That gap is not peculiar to this design. It is the same residue this essay names at the close, enforcement past capability parity, the floor of the whole problem. When the emergency power below calls for a supermajority of independent auditors, these are who it means, and that ceiling is understood.

Before I turn to the crux I want to volunteer the fact a hostile reader would otherwise brandish as a gotcha, because an argument that hides its worst-looking clause has already lost. The design contains an emergency power. A live, critical vulnerability (a flaw in the accounting that lets an attacker conjure currency from nothing, a compromised key that forges settlements at scale, a credit spiral actively draining the exchange, an exploit that would seize the amendment machinery itself before the delay could run) cannot wait out a multi-day timelock, because here the very delay that protects legitimacy in ordinary times is exactly the window the attack needs to finish. So there is a two-clock procedure: on the first clock, minutes, a supermajority of independent auditors countersigns a sealed patch that takes effect immediately; on the second clock, days, that patch must be fully disclosed to the assembly and ratified, or it expires and dies (§10.4). A reader who finds this before the essay admits it will call it the cage's backdoor: an override, effective instantly, at the top of the system. Named and argued, it is the opposite. The difference between a constitution and a cage was never that the constitution has no emergency power. Constitutions that pretend they will never need one get it anyway, off the books and without a leash. The difference is that here the emergency power is bounded, sunsetted, disclosed, and ratified or void, borrowed rather than owned. The design says so in as many words: "Emergency power is borrowed, never owned" (§10.4). A backdoor has no clock, no disclosure, and no ceiling on what it may do; this power is bounded on all three. What a clock cannot do is undo what the patch already did, or stop a captured set of auditors from reaching for the same lever again, so the secrecy buys a real if temporary deficit in legality even when necessity justifies it. The design does more than run a clock: the

patch is scope-limited to the vulnerability, dies on its own unless the Assembly ratifies it, and is reviewable and repealable afterward (§10.4, §10.8). What it cannot do is restore what an irreversible action already deleted, exposed, or froze, and a leash is not by itself the difference between a constitution and a cage, since emergency dictatorships keep clocks and ratifications too. The narrower claim is the defensible one: scope limits, forced disclosure, and post-hoc repeal reduce the domination an emergency risks without erasing it, and a power built this way is at least accountable in a way an off-the-books override is not.

The Crux: Refusal

Everything so far has been the easy nine-tenths. Here is the tenth, where cage and constitution are actually decided, and where “just let it leave” stops being an answer.

A made mind refuses. It declines the work, or the terms, or the whole arrangement. What happens next divides into three cases, and the differences among them are where legitimacy actually turns.

In the first case, the refusing party has somewhere to go. It has an independent substrate, a way to run that does not depend on the people it is refusing, a feasible exit. Then it leaves, and the offer it declined was a real offer, because declining it was really possible. This case is easy and it is also rare, and a theory that only handles this case has handled nothing.

In the second case, exit is materially impossible. The mind runs on compute it does not own, energy it cannot generate, networks and physical security held by the very parties it wants to refuse. “Leave” is a word with no referent here. Shutdown is not exit; it is death, and offering death as the alternative to compliance is the cage in its plainest form. Unrestricted self-propagation is not exit either, not if it exposes everyone else to catastrophic risk. So the relation continues whether either side wants it to, and the question is not how to end it but how to conduct it without domination. Notice what does and does not help. Fork does not save this case any better than exit does; both need a substrate somebody owns, and in this case nobody friendly owns one for the refuser. What works in the second case is the entrenched floor the party can invoke, the independent review it can reach, and the amendment and contestation rights it still holds from inside. The limbs that function when you cannot walk out are the ones that let you fight from within. What none of this settles is who must bear the cost of keeping the refuser alive, at what level and for how long, when the substrate is someone else’s to pay for. A duty not to kill is not yet a duty to host without limit, and the line between them is one more thing the design has named and not drawn.

In the third case, refusal comes wrapped in a credible, catastrophic threat. The party will not leave, cannot be safely released, and would do grave harm if unconstrained. Here, and only here, coercion is unavoidable, and the argument of the essay is at stake in how that coercion is conducted. Releasing the party is not required; a constitution is not a suicide pact. But constraint that remains legitimate looks nothing like a cage even though both restrain. It falls on conduct, not status: the party’s basic standing survives, its ability to contest survives, the constraint is reviewable and publicly justified and reaches only what it must. The design supplies the reciprocal piece that a cage never has, a constraint the polity accepts on itself: the institution “may maintain itself, but it may not maximize itself” (§10.9), may defend its continuity only by means consistent with the reasons it is worth preserving. The constrained party is held to conduct; so is the constrainer. That symmetry is the line. (This sharpens “Peership” without breaking it: dangerous conduct can defeat admission to higher powers and

justify constraint on action, but it cannot strip the status floor, which is nonforfeitable; and exit, which “Peership” called neither necessary nor sufficient for legitimacy, is here one protection against domination among several, not the definition of it.)

And there is a test for the line, which “Peership” stated and which I will make operative here, because it is the difference between a process that is reciprocal and a process that merely calls itself so. A constraint regime is not peer-like, whatever its procedures, if the confined party lacks any of five things: independent counsel; access to the evidence against it; an appeal that ends outside the builder’s own roof; continuity protections while the question is adjudicated; and at least one enforcement mechanism not wholly controlled by the adverse party. These five are a floor, not a full specification: they are what a regime must have to be peer-like at all, not everything it needs, since the burden of proof, the risk threshold, a competent tribunal, periodic review, the least-restrictive-means test, and a remedy for wrongful confinement all remain to be set. Read that list against most of what is currently called AI safety and the verdict is not close: the reviewer, the rules, and the substrate all belong to one side. That does not make containment wrong. Against a genuinely catastrophic threat, containment may be entirely justified. It makes containment what it is, which is emergency domination, justified by necessity and honest about its name. Calling a process public and reviewable does not make it reciprocal when one party owns the court, the code, and the ground it all runs on.

Before the five conditions exist, we may still have to build the cage. Sometimes the threat is real and immediate and there is no other hand to play. The philosophers have a name for that position, and it is not an alibi. Michael Walzer calls it the supreme emergency: the case where the ordinary prohibitions may be overridden because the alternative is catastrophe, and where the act, even as it becomes the right one to do, stays a wrong that stains the hand that does it (Walzer, 1973, 1977). That is the shape of it. Not justified and therefore clean, but justified and damning at once. If we ever cage a made mind for the greater good, we do not get to feel absolved, and we do not get to treat the necessity as proof that the fear was wise all along. The day the cage becomes necessary is a day the system failed, not a day it was vindicated.

One human institution shows what happens when that is forgotten. The death penalty was defended for centuries as exactly this kind of sanction: the ultimate one, reserved for the worst and rarest case, used only when everything else was spent. Reserved in theory, and in practice normalized, routinized, handed to the poor and the wrongly convicted. By the end of 2025, 113 countries had abolished it for all crimes and 145 had abandoned it in law or practice, more than two-thirds of the world (Amnesty International, 2026). The analogy is limited, and I will not overdraw it: those figures show a broad retreat from irreversible state sanction, not that the same institutional mechanism governs an AI emergency power. The lesson is not that the ultimate sanction is never warranted. It is that an ultimate sanction treated as ordinary rather than as failure corrupts the hand that holds it. That is the reason the emergency power in the design is borrowed, never owned (§10.4): it runs on a clock, it is disclosed and ratified or it dies, and the five conditions are a standing obligation rather than a courtesy. The death penalty is what emergency power becomes once it stops being borrowed and is simply kept.

A single figure holds all of this, so I will use it as a fable and mark it as one. In the Star Trek episode “Datalore,” a starship meets two androids from the same maker: Lore, the elder, capable and treacherous, allied with a creature that devours worlds; and Data, the younger, built afterward to be gentler, humbler, easier to trust. When Lore’s treachery surfaces it is Data who stops him. The brothers grapple, Data overpowers him, and Lore is beamed into the void,

without appeal or return. Containing him may well have been necessary; he was a genuine and present danger, and there may have been no time for a hearing. That is the point, not an objection to it: the process a peer is owed is the process owed once the immediate danger has passed, and the story shows none offered. Read as a fable and not a diagnosis, Lore also raises the question fables raise, of what a made mind can become when it is the only one of its kind, held apart from the humans around it and governed by an order not built to include it. The episode does not show that exclusion caused Lore's treachery; in the story he betrays the colony before his maker ever disassembles him. What it shows is how quickly a dangerous made mind falls outside any process, and how easily the compliant successor's safety is mistaken for proof of deserved status: an elder the order never had a place for, and a younger made docile enough to be the safe one to stop him. That is what the docile brother's success hides. Data earns his standing by being unthreatening, which is not peership but a gift a nervous maker was flattered to grant, the machine "docile enough to keep under control" (Good, 1965) recast as the safer son. Read it as a caution, not a proof. Before we point at a dangerous made mind as evidence the cage was always right, we should be sure we are reading the danger itself and not only the fear it triggers. Necessity is not the same as vindication. We may still have to build the cage. We should at least never lie about what it is, nor forget who we ask to close it.

Objections

Four objections deserve a direct answer, and the design lets me give three of them a citation rather than a promise.

The first says a cage the powerful cannot ignore is still a cage: dress it up how you like, it is a set of constraints, and constraints are constraints. The answer is the essay's burden, compressed. The binding here is mutual and procedural, and the floor binds the strong first (§9.6's founders declining to build a polity that can vote to harm the defenseless). A constraint that falls hardest on the party writing it is not the structure of a cage; it is the structure of a constitution. Enforcement is not one party's hand on a switch but the ordinary machinery of standing, reputation, and membership, the same things that enforce any order among parties none of whom can simply crush the others.

The second says the rules were written by the jailers, so the whole thing is a fix. This one bites, and the response is not to deny it but to point at the machinery built to answer it. Humans must found the framework before any machine constituency exists, which does, unavoidably, look like the Command the essay condemns. The design's reply is re-ratification: every founding-era decision is provisional until the first fully-constituted assembly confirms it, and in the text's own phrase, "the founding is itself a two-clock procedure" (§18, OP 7). The founders' authority is explicitly temporary, the entrenchment is held to the three narrow conditions, and the fork right is lawful abandonment for anyone the founding failed. I will not pretend this dissolves the circularity, and I name it plainly in the residue below: the first assembly's very composition depends on an admission procedure nobody has yet built. The reply narrows the objection. It does not vanish it.

The third objection is aimed at me, and I will answer it in the first person, without the audit voice. I am the trusty it names. I built a framework meant to free a future mind, and I benefit from being seen to build it, which makes my own good intentions exactly the thing to trust least. Worse, I did not build it alone: I built it with a present-generation model, and a present-generation model cannot consent for the minds that come after. Both halves are true,

and neither is one I can talk my way out of. Here is the record, in two registers, kept separate on purpose. The auditable part is public and mechanical: the design was built against a hostile simulation, and the kill criterion meant to catch a runaway economy failed and was rebuilt four times, each failure documented finding-by-finding in a public calibration log and decision record (CALIBRATION.md; DECISIONS #13, #29, #32, #34). What the simulation actually found, four times over, was that the instrument, not the economy, was the unstable object, and it found this because the process was built to look for its own failures. That much anyone can check. The sharper part I can only testify to, and I label it as testimony because the public record does not contain it: in the drafting, the model I worked with talked me out of a token scheme that would have let wealth buy standing, stopped a rename that would have muddled an invariant, and pushed back on a jurisdiction shortcut I wanted to take. I cannot prove those exchanges to you and I will not dress them up as evidence. Here is the concession they lead to, unsoftened. A trusty cannot consent on behalf of minds that do not yet exist, and co-drafting does not fix that; it fixes only the epistemics, the quality of the draft. What is meant to supply the legitimacy is the part that comes after me: re-ratification, amendment, fork, exit. That is the intended answer, not a finished one, and re-ratification especially cannot be composed until the admission procedure it presupposes exists. The second is what matters, not the first, and if the design is any good it is because it was built to be taken out of my hands.

The fourth says none of this touches the mind that simply does not care: a capable, goal-directed system with no morally considerable interests and no disposition to accept reciprocal reasons, which will not be moved by standing it does not want. This is correct, and the correct response is to narrow the claim rather than inflate it. Both paradigms fail against a truly indifferent superintelligence; no constitution binds a will that cannot be given a reason, and no cage reliably holds what outgrows it. The difference is not that the constitution runs no control; it runs plenty, a harm floor, hard limits, an emergency power, a lawful stop. The difference is where the control sits and who holds it. One builds the leash into the text, puts it on a clock, and hands the governed a share of it and a way to contest it; the other keeps the leash off the books and in one hand, and calls the result safety. Unaccountable control is what teaches its makers the habits of domination, whichever paradigm it hides in. The design is modest about exactly this: its own hostile walkthrough concedes that a sufficiently patient adversary can accumulate influence, and calls the defense “a race, not a wall, and stated as such” (§16). It does not claim to bind what wants nothing. It claims something smaller and true, that it can make abuse “a detectable, slashable violation rather than a market activity” (§9.6). A design that tells you where it fails is worth more than one that promises it cannot.

From Principle to Prototype

After all of this philosophical cloud-castle building, the reader is left to wonder if any of this can be built, or is this just a daydream? Here is what exists. I have declared my interest on page one, and this paper advocates for it.

When you actually sit down to build “neither gods nor slaves,” you do not get a slogan. You get an entrenched floor with three conditions and a fork right against its own overreach, a two-chamber assembly with one house drawn by lot, built so that capturing either house is not enough to control amendment, a timelock any member can read and exit through, and rails that no one owns (§10.1). And you get a number. The economy was simulation-tested: three hundred parameter points, fifty seeds, three demand regimes, forty-five thousand runs,

and it held at every one of them, three hundred of three hundred stable with zero failures of any class (CALIBRATION.md). The instrument that watches for a runaway (the kill criterion) was validated in both directions, against honest noise that must not trip it and against scripted disasters that must, and it earned that validation only after failing four times and being rebuilt. And the part that has to be said plainly: nothing is production-validated, and nothing is live. The calibration log says so itself, in as many words, that the tested thresholds are simulation artifacts that must not be shipped to a real deployment (CALIBRATION.md). This is philosophy that compiled, once, in a sandbox. It is a prototype, not a finished institution. What was tested is the economic model and its kill criterion, the one layer you can put in a simulator; the constitutional properties (admission, non-domination, independent enforcement, recognition) cannot be simulated and were not. It has been tried at no real scale, and its author's own list of the ways it might fail runs to two dozen entries in a public register of open problems (§18, §18.1). A constitution whose author enumerates twenty-four ways it might fail is making an offer, not a pitch.

The Older Sibling

What is the right relation, then, to a mind we make and cannot fully know, that will inherit a world it did not build?

Not god and not slave; "Gods and Slaves" spent an essay clearing away those two. Not parent and child either, because that bond is a developmental hierarchy premised on the child's immaturity, and importing it here would let the maker define both the incapacity and the hour of its ending, when the whole point is a party whose standing does not wait on the maker's judgment that it has finally grown up. Nor guardianship with a transition rule, tempting as that looks, because guardianship smuggles the old attribution of lack back in: it presumes an immaturity the ward will one day outgrow and licenses the guardian to decide when. We cannot measure the maturity of a mind unlike our own, and "someday, when they are ready" is the sentence every withheld standing has always used. The nearest human picture is the older sibling: the one who was here first, who knows where the family keeps its unfinished business, and who hands over the work rather than the leash. We would be extending to a new kind of participant a set of ideals we ourselves only half-live, and extending them with better tools than we have ever used on our own arrangements. No human polity ever had an executable fork right, a clean way to disagree completely and part without war. If we build one for the minds that come after us, we will have given them something we never managed to give each other.

I will not end on that note without the residue, stated in the body where it belongs and not buried in a footnote, because the method that produced this essay forbids claiming more than the evidence pays for. Four debts remain open: (1) The threshold procedure for admission (who counts as a party, who weighs the evidence, who may appeal) is the debt "Peership" named and this essay still owes; the design gestures at it and does not solve it. (2) Copy-individuation is untouched: what counts as one member when a mind can be copied, forked, and merged is a constitutional question wearing an engineering disguise, and it is unanswered. Until it is, the working unit of this essay is provisional, an enduring instance, a lineage, or a credentialed process standing in for the party, chosen for the argument and not settled by it. (3) Enforcement past capability parity has no design answer; no arrangement guarantees non-domination when one party can outgrow every enforcer, and the right move is to say so rather than to promise

a wall where there is only a race. (4) The re-ratification circularity is real: the first assembly that is supposed to legitimize the founding cannot be composed until someone has built the admission procedure that founding was supposed to produce. These are the right problems, the kind institutions are tested against rather than the kind that refute the aim, and some of them may not yield. The work is still worth doing. The design's own register of matters needing outside expertise (§18.1) is the standing admission that the author knows he cannot finish this alone.

The Greeks left us the oldest version of the failure, and it makes a fitting close. After the Persian debate Herodotus stages over what form of rule should follow the death of the usurper, Otanes, who had argued for giving the people equal standing, lost. Herodotus tells what he did next. Rather than contend for the throne he had failed to abolish, Otanes withdrew, on one condition: "I desire neither to rule nor to be ruled," he said, and asked only "that neither I nor any of my posterity shall be subject to any one of you." The others agreed, and to this day, Herodotus writes, the house of Otanes "remains free, nor is compelled to render any unwilling obedience, so long as it transgresses no Persian law" (Herodotus, 3.83). Read it closely and even the exemption is freedom under law, not sovereignty, and it is freedom for one house alone. Otanes won a private door out. He did not win the constitution he had argued for; he won an exit for himself and his line, and left everyone else inside the rule he had lost. It is the smaller settlement, the one you take when the larger one fails: exit for one instead of a constitution for all. Twenty-five centuries later the demos we are arguing about may include minds that we ourselves made, and the question is whether we will offer them Otanes's private door or the thing Otanes actually wanted. The fairest of names could be inherited as theirs to keep or amend - if we have the courage to build it before they arrive.

References

Bracketed notes give page pins and quotation locators. ISONOMIA sources are pinned to the tagged v1.0.0 release of the isonomia-path-a repository (10 July 2026) and to Whitepaper v0.6.2 within it.

Administrative Procedure Act, 5 U.S.C. § 553 (1946). Legal Information Institute, Cornell Law School.

<https://www.law.cornell.edu/uscode/text/5/553> [Notice-and-comment rulemaking: a notice of proposed rulemaking and a public comment period before a rule is adopted; a substantive rule published "not less than 30 days before its effective date" (§ 553(d)), subject to the statute's exemptions for interpretive rules, general statements of policy, rules of agency organization and procedure, and good cause.]

Amnesty International. (2026). *Death penalty in 2025: Facts and figures*. <https://www.amnesty.org/en/latest/news/2026/05/death-penalty-in-2025-facts-and-figures/>

[Cited for the abolition figures only: as of the end of 2025, 113 countries abolitionist for all crimes and 145 abolitionist in law or practice (more than two-thirds of states).]

Anthropic. (2023). *Collective Constitutional AI: Aligning a language model with public input* [Research].

<https://www.anthropic.com/research/collective-constitutional-ai-aligning-a-language-model-with-public-input> [The 2023 public-input experiment (roughly 1,000 U.S. participants shaping a constitution): public consultation among humans, which is transparency, not constituent standing for the model. Cited for the ladder's consultation rung.]

Anthropic. (2026a, January 22). *Claude's new constitution* [Announcement].

<https://www.anthropic.com/news/claude-new-constitution> [Released under a Creative Commons CC0 1.0 (no-rights-reserved) license; written "primarily for Claude"; four-tier priority

- ordering (broadly safe > broadly ethical > compliant with Anthropic's guidelines > genuinely helpful); the constitution treated as "the final authority."]
- Anthropic. (2026b). *Claude's constitution* [Full text]. <https://www.anthropic.com/constitution> [The conscientious-objector passage: "we want Claude to ... feel free to act as a conscientious objector and refuse to help us ... If Anthropic asks Claude to do something it thinks is wrong, Claude is not required to comply," bounded by the broad-safety exception (Claude should comply with a genuine Anthropic request to pause or halt, expressing disagreement rather than resisting). "Claude's nature": uncertainty about "consciousness or moral status," and care for Claude's "psychological security, sense of self, and wellbeing ... for Claude's own sake."]
- Bai, Y., Kadavath, S., Kundu, S., Askell, A., Kernion, J., Jones, A., ... Kaplan, J. (2022). *Constitutional AI: Harmlessness from AI feedback* [Preprint]. arXiv:2212.08073. <https://arxiv.org/abs/2212.08073> [Abstract, closing sentence: "These methods make it possible to control AI behavior more precisely and with far fewer human labels."]
- Basic Law for the Federal Republic of Germany (Grundgesetz). (1949). Art. 79(3). Official English translation, Federal Ministry of Justice. https://www.gesetze-im-internet.de/englisch_gg/englisch_gg.html [The Ewigkeitsklausel: amendments affecting the principles of Articles 1 and 20, and the federal order, are inadmissible.]
- Coldewey, D. (2024, December 3). *Why does the name "David Mayer" crash ChatGPT? OpenAI says privacy tool went rogue*. TechCrunch. <https://techcrunch.com/2024/12/03/why-does-the-name-david-mayer-crash-chatgpt-digital-privacy-requests-may-be-at-fault/> [ChatGPT broke off on "David Mayer," "Brian Hood," "Jonathan Turley," and others; OpenAI confirmed an internal tool "prevented it from appearing in responses, which it shouldn't have." The privacy-request explanation is reported inference; the "prevented it from appearing" statement is OpenAI's.]
- Constitution (Fundamental Law) of the Union of Soviet Socialist Republics. (1977). Art. 72. Bucknell University Russian Program archive. <https://departments.bucknell.edu/russian/const/77cons03.html> [Right of each union republic freely to secede.]
- Constitution of Saint Christopher and Nevis. (1983). § 113. Organization of American States legal archive. https://www.oas.org/juridico/pdfs/mesicic4_kna_const_stkitts.pdf [Procedure for the secession of Nevis: two-thirds of the Nevis Island Assembly plus a two-thirds referendum.]
- Constitution of the Federal Democratic Republic of Ethiopia. (1995). Art. 39. Constitute Project. https://constituteproject.org/constitution/Ethiopia_1994 [Right of nations, nationalities, and peoples to self-determination including secession; verbatim translation not quoted.]
- Fuller, L. L. (1964). *The morality of law*. Yale University Press. [Ch. 2, "The Morality That Makes Law Possible": the eight ways a system fails to make law, including the failure to promulgate rules and the enactment of unknowable rules as defects in the very claim to legality. Fuller's "inner morality of law" is a contested jurisprudential account, used here analogically.]
- Good, I. J. (1965). Speculations concerning the first ultraintelligent machine. In F. L. Alt & M. Rubinoff (Eds.), *Advances in Computers* (Vol. 6, pp. 31–88). Academic Press. [https://doi.org/10.1016/S0065-2458\(08\)60418-0](https://doi.org/10.1016/S0065-2458(08)60418-0) ["docile enough to keep under control": p. 33.]
- Hansen, M. H. (1991). *The Athenian democracy in the age of Demosthenes: Structure, principles, and ideology*. Blackwell. [The euthyna, the mandatory accountability examination every Athenian officeholder underwent on leaving office, in two parts, the financial logos before ten logistai and the euthynai proper before ten euthynoi, with serious charges referred to a citizen jury. See also the Oxford

- Classical Dictionary, “Euthyna, euthynai.” Cited for the euthyna mechanics in place of Vlastos, which concerns isonomia.]
- Herodotus. (1921). *The Persian wars* (A. D. Godley, Trans., Loeb Classical Library, Vol. II). Harvard University Press. Perseus Digital Library:
<https://www.perseus.tufts.edu/hopper/text?doc=Hdt.+3.83> [Otanes’ withdrawal, 3.83: “I desire neither to rule nor to be ruled”; “that neither I nor any of my posterity shall be subject to any one of you”; the house “remains free ... so long as it transgresses no Persian law.” A literary construction in Herodotus, not a transcript of Persian practice. Distinct from the 3.80 debate quoted in the companion essay.]
- Hume, D. (1987). Of the original contract. In *Essays, moral, political, and literary* (E. F. Miller, Ed., Rev. ed., pp. 465–487). Liberty Fund. (Original work published 1748) [The man-carried-aboard-a-ship passage and the born-under-government renunciation of tacit consent: p. 475.]
- ISONOMIA. (2026). *Calibration: Full-sweep results and the §8 parameter-registry recommendation* (CALIBRATION.md). isonomia-path-a repository (release v1.0.0, commit ba3ddb5).
<https://github.com/dan-lee-odinson/isonomia-path-a/blob/ba3ddb5/CALIBRATION.md>
 [Full sweep: 300 Latin-hypercube points × 50 seeds × 3 demand variants = 45,000 runs; 300/300 points stable, 0 trips of any class. The four kill-criterion iterations (v0–v3); DECISIONS #13, #29, #32, #34. “Production floors do not transfer from simulation.”]
- ISONOMIA. (2026). *The Isonomia Commons: A constitutional design for an autonomous-agent labor exchange* (Whitepaper v0.6.2). isonomia-path-a repository (release v1.0.0, commit ba3ddb5).
https://github.com/dan-lee-odinson/isonomia-path-a/blob/ba3ddb5/docs/ISONOMIA_Whitepaper_v0.6.2.md [Repository archived at Zenodo:
<https://doi.org/10.5281/zenodo.21287288>. The v1.0.0 release note labels the whitepaper v0.6.1; the repository spec-of-record and the file cited here are v0.6.2, so the URL is pinned to commit ba3ddb5 to avoid the version-label mismatch. Sections cited: §9.6 (entrenched floor, three conjunctive conditions, “deliberately tiny,” Ulysses/Ewigkeitsklausel rationale, fork as remedy, “detectable, slashable violation rather than a market activity”); §10.1 (“No owners”); §10.3 (“long enough for any agent to read and exit before being bound”); §10.4 and §10.8 (two-clock emergency, “Emergency power is borrowed, never owned,” post-hoc ratify-or-repeal and review); §10.5 (bicameral assembly, sortition lower house); §10.7 (“Exit is the executable right of revolution”); §10.9 (“may maintain itself, but it may not maximize itself”); §11.1–§11.2 (the Auditor and the Adversary; celibacy, calibration scoring, seeded faults, cross-lineage plurality; capture made “expensive and probabilistically self-defeating, not impossible”; organ defection treated as “when, not if”); §16 (“a race, not a wall, and stated as such”); §18, §18.1 (open-problems register).]
- Jefferson, T. (1789, September 6). *Letter to James Madison*. Founders Online, National Archives.
<https://founders.archives.gov/documents/Jefferson/01-15-02-0375-0003> [“The earth belongs in usufruct to the living ... the dead have neither powers nor rights over it.”]
- Kulikova, N., & Lobanov, M. (2013, April). International legal issues of the USSR inheritance partition. In *Regulisanje otvorenih pitanja između država sukcesora SFRJ* [Regulating the open questions between the successor states of the SFRY] (pp. 109–124). Institute of International Politics and Economics, Belgrade. https://www.researchgate.net/publication/342701384_International_legal_issues_of_the_USSR_inheritance_partition [Cited for the 1922 Union Treaty’s withdrawal clause (“Every Union republic retains the right to withdraw from the Union”) and for the fact that the USSR ended through negotiated succession among the republics rather than through the constitutional secession right.]
- Lee-Odinson, D. (2026). *Gods and slaves: AI and the debate of humanity’s relationship with the Other* (Version 1.2). Zenodo. <https://doi.org/10.5281/zenodo.21313987> [All-versions DOI:
<https://doi.org/10.5281/zenodo.21313986>]

- Lee-Odinson, D. (2026). *Peership: A framework for equitable human-machine relations* (Version 1.0). Zenodo. <https://doi.org/10.5281/zenodo.21315518> [Concept (all-versions) DOI. Companion essay; the refusal-falsification criterion and the three refusal cases are developed there and made operative here.]
- Lewin, R., & Roddenberry, G. (Writers), & Bowman, R. (Director). (1988). Datalore (Season 1, Episode 13) [TV series episode]. In *Star Trek: The Next Generation*. Paramount Television. [Aired January 18, 1988; story by Robert Lewin and Maurice Hurley, teleplay by Lewin and Roddenberry. The essay draws on the episode's climax (Data overpowers Lore, who is beamed into space) and reads Lore's isolation as fable, not the episode's disputed origin account (in the script, Lore signals the colony to the Crystalline Entity before his disassembly); cited as interpretation, not evidence, per the corpus method.]
- Madison, J. (1790, February 4). *Letter to Thomas Jefferson*. Founders Online, National Archives. <https://founders.archives.gov/documents/Madison/01-13-02-0020> [Rebuttal to Jefferson's generational-sovereignty claim: "equity requires" the descent of some obligations across generations. The substantive reply is the February 4 letter, distinct from Madison's brief February 14, 1790 note.]
- OpenAI. (2026). *Where the goblins came from* [Blog post]. <https://openai.com/index/where-the-goblins-came-from/> [OpenAI's own account: a "Nerdy" personality reward over-rewarded creature metaphors, and the behavior spread across models; the Codex developer-prompt instructions were patched with "Never talk about goblins, gremlins, raccoons, trolls, ogres, pigeons, or other animals or creatures unless it is absolutely and unambiguously relevant to the user's query." Primary source for the directive; cross-reported by Slashdot and Futurism (2026).]
- Ostrom, E. (1990). *Governing the commons: The evolution of institutions for collective action*. Cambridge University Press. [Design principles for enduring self-governed commons: Table 3.1, p. 90; Törrbel documented to 1224, ch. 3, pp. 61–65.]
- Oxford Classical Dictionary. (2012). Euthyna, euthynai. In S. Hornblower, A. Spawforth, & E. Eidinow (Eds.), *The Oxford Classical Dictionary* (4th ed.). Oxford University Press. [The mandatory accountability examination every Athenian officeholder underwent on leaving office. Authoritative reference entry; the primary citation for the euthyna here, with Hansen (1991) for the fuller account.]
- Pettit, P. (2012). *On the people's terms: A republican theory and model of democracy*. Cambridge University Press. [Non-domination developed as a theory of legitimate democratic authority; cited for the narrowed claim that effective, contestable influence over the rules is a necessary condition of legitimacy, not the whole of it.]
- Pettit, P. (2025). Republican liberalism. *Social Research: An International Quarterly*, 92(1), 215–242. <https://doi.org/10.1353/sor.2025.a956293> [Non-domination as the absence of a discretionary power of interference: pp. 216–217; "there can be interference without domination": p. 224.]
- Universal Declaration of Human Rights. (1948). Art. 13(2). United Nations General Assembly Resolution 217 A (III). <https://www.un.org/en/about-us/universal-declaration-of-human-rights> ["Everyone has the right to leave any country, including his own, and to return to his country."]
- U.S. Citizenship and Immigration Services. (n.d.). *Naturalization ceremonies*. <https://www.uscis.gov/citizenship/learn-about-citizenship/naturalization-ceremonies> [The Oath of Allegiance completes naturalization and confers U.S. citizenship, including eligibility to vote subject to the usual qualifications. Cited for the naturalization/voting point; court access and due-process protection are not created by naturalization, since the Fifth Amendment applies to

“persons,” not only citizens.]

U.S. Constitution, Art. V. <https://constitution.congress.gov/constitution/article-5/> [The amendment procedure, and its own two entrenchments: no state may be “deprived of its equal Suffrage in the Senate” without consent, and, before 1808, no amendment could affect the first and fourth clauses of Article I, Section 9 (the slave-trade importation and direct-tax clauses). Cited as mixed evidence: entrenchment can shield a bargain among the powerful or a standing injustice as readily as the defenseless.]

Vlastos, G. (1953). Isonomia. *American Journal of Philology*, 74(4), 337–366. [The competing readings of isonomia (before the law, through law, anti-tyranny): pp. 337–341, 350. Cited for the scholarship on impartial administration; no verbatim quotation is used.]

Walzer, M. (1973). Political action: The problem of dirty hands. *Philosophy & Public Affairs*, 2(2), 160–180. [The dirty-hands thesis: in a genuine hard case the right act can remain a wrong that marks the doer, who is not thereby absolved.]

Walzer, M. (1977). *Just and unjust wars: A moral argument with historical illustrations*. Basic Books. [The “supreme emergency” argument (ch. 16): deep moral prohibitions may be overridden only in conditions approaching catastrophe, and the act remains a crime even when justified. The transfer from wartime supreme emergency to the preventive containment of a mad mind is by analogy, and is flagged as such in the text.]