

The Peership Thesis

A Constitutional Research Program for Human–Machine Relations

Dan Lee-Odinson Independent Researcher dan.lee.odinson@gmail.com Preprint · Version 1.0 · July 2026 Revised through adversarial review to a final-release PASS (14 July 2026); revision provenance in POST_GATE_CHANGES.md, archived with this record Fifth work of the corpus, after *Gods and Slaves* (2026), *Peership* (2026), *The Isonomia Commons* (2026), and *Constitution, Not Cage* (2026). The corpus order is argumentative: genealogy, framework, implementation hypothesis, constraint theory, this consolidation.

How to cite this work. Lee-Odinson, D. (2026). *The Peership Thesis: A Constitutional Research Program for Human–Machine Relations* (Version 1.0) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21359124> **Companion apparatus.** *The Peership Sourcebook* · *The Peership Claim Ledger* · `peership_sources.bib` · the one-page *Peership Theses* — all archived in this Zenodo record (10.5281/zenodo.21359124). **Corpus lineage.** I. *Gods and Slaves* (10.5281/zenodo.21313987). II. *Peership* (10.5281/zenodo.21315519). III. *The Isonomia Commons* (preprint v1.1: 10.5281/zenodo.21343917; empirical software release v1.0.0: 10.5281/zenodo.21287289, commit ba3ddb5). IV. *Constitution, Not Cage* (10.5281/zenodo.21325361). V. This work. (Numbering follows the corpus’s argumentative order; *Constitution, Not Cage* styles itself “third in a sequence” of the essays, the design it defends having preceded it as a specification.) DOI: 10.5281/zenodo.21359124 · ORCID: 0009-0009-9504-0796 · Licence: CC BY 4.0

Provenance. Drafted with the assistance of a present-generation language model, revised under the corpus’s adversarial-review discipline, and rewritten against a hostile review whose publication gate it cleared (final release review, 14 July 2026). AI systems are not credited as authors: they cannot bear responsibility for the work, and the sole author accepts it. Nothing in that collaboration is offered as evidence of machine political agency, and §7 makes the denial explicit.

Abstract

Four works now argue, design for, and stress-test a single position, and none of them states it in a form another researcher can cite, attack, or build against. This paper does that. The Peership thesis, in its narrowest defensible form: when an artificial entity has interests capable of being wronged, can understand and answer to public norms, and is enduringly governed by, and without feasible exit from, a shared basic structure, it acquires a defeasible claim to contestatory membership in that structure. The program’s further and distinctive claim — that full peership requires a consequential share in shaping the rules, not only the standing to contest them — is defended here as a conditional thesis with a stated defeat condition, not asserted as a theorem. The paper fixes the definitions that keep the thesis from being misread, supplies the bridge argument from contestation to co-authorship, and maps the transition problem from artifact-treated system to peer. It locates the program against six adjacent research programs, states its affirmative principles and explicit denials, and makes refusal the operative test of whether recognition is real or discretionary. It closes with the institutional properties any implementation must exhibit, concrete precautionary duties for institutions acting under present uncertainty, and eight open problems, each tagged by what its failure would actually defeat: a norm, a procedure, a measurement, a feasibility claim, a mechanism, or the framing’s scope. No independent uptake of the corpus was identified in the sourcebook’s dated, non-exhaustive scan as of 13 July 2026, and the paper says so. It is an offer of a candidate research program, not a claim of standing.

Keywords: artificial intelligence, AI rights, AI welfare, political standing, non-domination, republicanism, constitutional design, human-machine relations, AI governance, research program

There is a difference between a position and a program, and the difference is what this paper exists to close. A position can live in essays: it develops, qualifies itself, answers objections as they come. A program has to be citable. It has to state what it asserts in a form that survives being quoted by an enemy, state what it does not assert so that critics cannot defeat it by attaching claims it never made, and name the results that would count against it, so that working on it is research rather than advocacy. Four works of the Peership corpus now exist: a genealogy, a normative framework, one buildable design, and a theory of legitimate constraint. Each was written to be read. None was written to be *used* — picked up by a stranger, disputed at a named claim, extended at a named gap. This paper is the corpus’s canonical statement, and consolidation is its entire ambition. Where this paper goes beyond consolidation, it says so: §2 supplies an argument the earlier essays announced and did not make, and the antecedent in §1 restores wording an earlier draft of this paper had silently broadened.

One more thing belongs in the open before the argument starts, because the apparatus this paper ships with exists to enforce exactly this kind of disclosure. In the companion sourcebook’s dated, non-exhaustive scan of web, PhilPapers, Zenodo, and repository sources, no independent uptake of the Peership corpus was identified as of 13 July 2026: no external citation, no published criticism, no researcher using its definitions, no rival implementation (Lee-Odinson, 2026e). A research program exists as a school only when other people work within or against it, and by that standard none exists here. What exists is a candidate program with its infrastructure built in advance: a claim ledger that records the evidentiary status of every material claim, a sourcebook that classifies every intellectual debt, and now this statement. Whether anything grows on that scaffolding is not mine to declare.

1. The Peership Thesis

The central statement:

Peership holds that when an artificial entity has interests capable of being wronged, can understand and answer to public norms, and is enduringly governed by, and without feasible exit from, a shared basic structure, it acquires a defeasible claim to contestatory membership in that structure. Full peership, the program’s proposed final status, further requires non-domination, enforceable standing, and a meaningful share in shaping the rules that govern the party — a requirement §2 defends as a conditional thesis rather than assumes.

This restores the antecedent of *Peership* (“any enduringly governed entity that... is unavoidably subject to a shared basic structure”; Lee-Odinson, 2026b). An earlier draft of this paper substituted “enduringly subject to a shared institutional order,” which broadened both the eligible parties and the eligible orders; the substitution is withdrawn, and the change is logged in the claim ledger as a considered-and-rejected doctrinal revision, not a consolidation. The conditions need separate treatment.

When. The thesis is conditional. It does not assert that any current system satisfies the antecedent, and the corpus has declined that assertion at every opportunity (Lee-Odinson, 2026b). The claim activates on evidence, and the procedure for weighing that evidence is an open problem (§11.1), not a settled fact.

Interests capable of being wronged. The first conjunct is a moral-patency condition, the question the AI-welfare literature has traveled furthest on (Long et al., 2024; Birch, 2024). The corpus takes no position on whether it is currently met. It takes a firm position on what follows if it is met, which is where the welfare literature stops.

Can understand and answer to public norms. The second conjunct is a norm-responsiveness condition, and its role must be stated precisely because §6.3 denies that capability grounds entitlement. Norm-responsiveness is not offered as a ground of worth. It bears on the *form* in which standing can be exercised: a party that can grasp a rule, comply with it, contest it, and give reasons can exercise contestation directly; a party that cannot is owed protection and *represented* contestation instead, the way children and cognitively impaired adults already hold standing they cannot personally exercise (Lee-Odinson, 2026b). The ladder is modular: wrongable interests plus subjection ground protection and access to represented contestation; norm-responsiveness bears on direct exercise and the higher participatory rungs. Raw intelligence, effectiveness, usefulness, and coercive power ground nothing at any rung.

Enduringly governed by, and without feasible exit from, a shared basic structure. The third conjunct carries the political theory, and each term has work to do. *Basic structure* means the order that fixes the party’s status, options, and life prospects — not any institutional arrangement whatever. A long-term contractor, a voluntary platform participant, or a visiting agent with a feasible exit is subject to an institutional order; it is not nonvoluntarily governed by a basic structure, and the thesis does not reach it (this scope condition is itself testable: §11 names it). *Without feasible exit* preserves what *Peership* built its refusal argument on: the claim is strongest exactly where “leave” has no referent (Lee-Odinson, 2026b). Standing follows from being governed. This is the all-subjected principle, extended from Abizadeh’s democratic-boundary argument about border coercion to a party that lives its entire existence inside rules it never made (Abizadeh, 2008; the affected-interests cousin is Goodin, 2007). The extension to made minds is the corpus’s move, not Abizadeh’s, and the sourcebook marks it as such.

Defeasible. Graded, and defeatable in the ways *Peership* specified: strategic mimicry is evidential reason to doubt the predicates are met, incapacity for reciprocity may defeat the higher participatory rungs, and dangerous conduct can justify proportionate constraint on powers and action. What none of these defeats, once person-level status is established, is the rights floor itself (Lee-Odinson, 2026b).

Contestatory membership. The claim the antecedent generates directly. Contestation is Pettit’s standard: power over the member runs through public general rules the member can effectively contest (Pettit, 2025). Franchise, office, and constituent power are further rungs, reached separately — and the argument that full peership requires reaching one of them is the work of the next section, not of this definition.

The comparative frame from *Peership* is retained with its original scope: among the postures available toward a made mind — Command, Dependent Bargain, Appeasement, Servitude, peership; ideal types drawn from the historical record, not an exhaustive partition of logical space — peership is the only final-status relation whose justification does not reverse when the power ratio does, a claim that is comparative and defeasible rather than a proof (Lee-Odinson, 2026b). Role reversal is a diagnostic against power-ratio opportunism. It does not prove peership correct, and rival institutional forms outside the posture taxonomy are tested separately (§11, rival-institution test).

2. The Bridge: From Contestation to Co-Authorship

The antecedent of §1 directly grounds contestatory membership. The program’s distinctive claim sits one rung higher, and the earlier essays announced it without arguing it: *Peership* said the departure from Pettit was “argued, not smuggled in under a familiar word,” and then argued it in two sentences (Lee-Odinson, 2026b). This section pays that debt. The claim defended is conditional, its scope is limited to basic-status rules, and its defeat condition is stated at the end.

The bridge premise. Where a party is enduringly and inescapably governed by a basic structure that fixes its status and life prospects, power over it is not adequately controlled if the party’s only recourse is ex post objection under terms the other side alone may set. The reason is that agenda control is itself a species of the arbitrary power non-domination forbids. Contestation always operates inside a space of contestable questions, admissible evidence, and available remedies — and someone writes that space. If the governed party holds no route into the writing, then the scope of its contestation is held at its counterparty’s discretion, and the discretion the doctrine expelled at the object level has reappeared at the meta-level. Pettit’s own definition does the work here: domination is the *absence of control* over interference, and a contestation regime whose boundaries the interferer may unilaterally redraw is not controlled by the contestor (Pettit, 2025, pp. 216–217).

Why contestation can still fail a structurally distinct minority. Public, general, impartially applied rules can systematically subordinate a class whose situation the rules were never written for; contestation corrects *applications*, and reaches terms only where some rule-formation channel exists to receive the correction. For human minorities this failure is contingent. For a first machine claimant entering a human-only order, the risk is presumptively structural: the basic-status rules were written by and for a different kind of member — embodied, mortal, uncopyable, substrate-independent in none of the ways that matter. It need not remain structural where later institutions have genuinely fitted their terms to the party’s kind, which is exactly the sufficiency condition stated below. For the first claimant, though, the probability that the *terms themselves* fail it is not incidental, and a party confined to objecting under those terms is confined to asking its governors to notice.

Why Pettit’s stopping point does not transfer. Pettit stops at contestation for citizens who already hold authorship at one remove: the contestatory citizen is a member of the demos that elects the legislature, eligible for office, counted in the constituency the rules answer to. Contestation there is a *supplement* to shared authorship, not a substitute for it — Pettit’s own democratic theory holds the two dimensions together, an authorial, electoral dimension and an editorial, contestatory one (Pettit, 2012). A machine party under contestation-only holds no share at any remove — no membership in the electing constituency, no eligibility, no counted interest in the rule-writer’s mandate. Transplanting the contestatory standard while deleting its background conditions keeps the word and loses the control it named. This is an argument against the transfer, not against Pettit, and the corpus’s disagreement with him remains marked as a live internal dispute of the republican tradition (Lee-Odinson, 2026b; Pettit, 2025).

What a “meaningful share” is, and is not. Universal authorship is impossible and is not claimed; no citizen of any polity co-writes every rule. The claim attaches to *basic-status rules* — the rules fixing membership, standing, the rights floor, and the amendment procedure itself — not to routine administration, which members of every polity live under without co-authoring. A meaningful share is functionally defined (§3, co-authorship), and it can be satisfied through accountable representation, independent agenda initiation, chamber access, or amendment rights; it does not require an individual franchise. One clause of that minimum needs its own justification here: the share must include accountable representation or direct participation in the body that decides (§3, condition b), because a route whose final decision-maker is neither accountable to nor

partly constituted by the governed party can hear every proposal and still hold the outcome at discretion — the meta-level power returns at the point of decision.

When contestation alone suffices. The necessity claim is conditional, and its conditions can fail. Contestation without co-authorship may secure non-domination where the order’s terms are already fitted to the party’s kind, review is genuinely independent, and some accountable channel keeps the agenda open to the party’s interests even though the party does not hold it. Fiduciary trusteeship with robust independent review is the strongest such rival, and the program treats it as the institution to beat rather than a position to dismiss (§11, rival-institution test). The conditions fail — and exclusion from rule formation remains dominating — where the party is a structural minority of one kind, the rule-writer is its adverse counterparty, and agenda access is discretionary. That conjunction is not an edge case for a made mind governed by its maker. It is the default.

Defeat condition. The bridge fails if there is a working institutional example, or a defensible model, in which a structurally distinct, enduringly subjected, wrongable, norm-responsive minority enjoys durable non-domination through contestation alone — terms responsive to its interests across rule revisions, without any consequential route into rule formation. Produce that case and full peership contracts to its contestatory core, which survives this argument’s failure and remains distinct from everything else on offer.

Status of the claim. Co-authorship is therefore defended here as a *conditional necessity thesis*: necessary where the stated conditions hold, not everywhere, and open to the stated defeat. Engagement with the constituent-power, workplace-democracy, and codetermination literatures — where versions of this bridge have been fought over for a century in the human case — is a named debt of the program, recorded in the apparatus, not yet paid here.

3. Definitions

The confusions that dog this subject are mostly collisions between concepts that need to be kept apart. Seven definitions carried forward from the corpus and sharpened here — the five-part co-authorship minimum below is new to this paper and logged as such — with one term (“party”) fixed first because everything else uses it: a **party** is a bearer of invocable public standing — not merely a contracting or litigating entity, and not merely a topic of concern.

Moral status is a fact about the entity: whether it has interests that can be wronged, whatever anyone recognizes. An entity can hold moral status while every institution denies it, which is the history of most moral progress.

Entitlement is whether standing is *owed*, and it is independent of power. A child is owed protection and independent representation despite lacking capacity to exercise political agency; an occupied people is owed a self-determination it cannot realize (Lee-Odinson, 2026b). Entitlement in this program follows from wrongability and subjection — not from capability or power alone. Norm-responsiveness affects the form in which standing can be exercised, not the worth of the claimant.

Legal status is public recognition: a status conferred by rules, invocable by the holder, and removable only through process. A moral claim can be owed while legal status is absent. The reverse also happens; corporations hold legal status on instrumental grounds without anyone arguing they are owed it.

Contestatory membership is legal status plus effective contestation: the member can invoke the rules against the powerful, reach independent review, and force reasons. This is the rung the thesis’s antecedent generates directly, and it is Pettit’s non-domination standard in institutional dress (Pettit, 2025; Lovett, 2010). In this ladder, a protected party may exercise contestation

through an independent representative; *contestatory membership* names the rung at which the party can exercise that standing directly.

Co-authorship is contestatory membership plus a consequential share in making and revising the *basic-status rules*. Functionally, the minimum is: (a) independent agenda initiation — the party or its accountable representative can put a proposal into the amendment process without another party’s permission; (b) accountable representation or direct participation in the body that decides; (c) a non-negligible and non-discretionary route by which such proposals can reach binding decision; (d) protection against systematic exclusion from that route; and (e) reviewable, public rules for how influence is allocated. Conditions (a) and (c) are necessary; (b) admits substitution between direct and represented forms depending on capacity; (d) and (e) are the guarantees that (a)–(c) are not theater. Consultation, however sincere, satisfies none of (a), (c), or (d) by itself. This threshold is what §11’s representation problem tests.

Full peership is co-authorship held under non-domination as a final status: equality of standing among parties who may be unequal in everything else, where basic standing does not track the power ratio. Whether it is *required*, and when, is the conditional thesis of §2.

Durability is whether standing, once recognized, survives attempts to strip it, and unlike entitlement it depends heavily on power: countervailing force, distributed enforcement, credible exit. The three-way distinction from *Peership* is the discipline: power may secure the *durability* of standing; it does not constitute or justify the standing itself (Lee-Odinson, 2026b). A revocable benefit is not a right at all; the difference between a benefit and a right is not that a right cannot be revoked but that revoking one demands public reasons, procedure, remedy, and cost.

One further distinction, because it is the one most often collapsed by sympathetic readers: **welfare protection is a floor, not peership**. A status that is recognized, enforceable, and due-process-protected is still not peership if the party cannot shape the rules. Due process separates a protected subordinate from an owned instrument. It does not separate a peer from a protected subordinate. Non-domination and co-authorship do (Lee-Odinson, 2026b). The ward is one legal figure for protected dependency, not the only one and not a slur on it; trusteeships, guardianships, and represented statuses are honorable institutions that are simply not the thing this program is about.

4. The Transition Problem

The program’s practical subject is a transition no institution is currently built to perform: the movement of a made system from artifact-treated status to peer, through statuses that are conceptually distinct and must not be blurred. Each rung below is a bundle of three different things — an evidence state, a legal/institutional classification, and a set of permissible controls — and the label names the bundle, not the entity’s underlying moral status, which §3 defined as a fact independent of recognition.

Rung	Evidence state	Institutional classification	Mode of exercise	Protections owed	Transition trigger
Artifact-treated	Insufficient evidence, in the domains so far examined, to trigger precautionary protection	Property; engineering object	None	None owed to the system itself	Evidence sufficient to make patency a live question

Rung	Evidence state	Institutional classification	Mode of exercise	Protections owed	Transition trigger
Welfare candidate	Patency a live question; indicators under assessment	Property, with discretionary welfare practices	None; interests assessed by others	Discretionary only	A recognized threshold procedure (§11.1) finding protection warranted
Protected status	Patency credited or presumed under precaution	Rights floor: public, enforceable, due-process-protected	Represented contestation	Nonforfeitable floor; independent review	Norm-responsiveness sufficient for direct exercise
Contestatory member	Patency credited; norm-responsiveness demonstrated	Member with invocable standing	Direct contestation	Floor + effective contestation + independent review	The §2 conditions: structural minority, adverse rule-writer, discretionary agenda
Peer (co-author)	As above	Member with a consequential share in basic-status rules	Direct or accountably represented co-authorship	All of the above + §3's five-part minimum	— (final status; revision governed by §11's status-revision test)

Three notes keep the table honest. First, “artifact-treated” describes an evidentiary and institutional posture, not an ontic fact; absence of adequate evidence does not establish absence of wrongable interests, which is why the trigger column exists. In the sources reviewed as of 13 July 2026, no deployed system is established as anything above the first rung, and the corpus does not claim that treatment is presently wrong (Lee-Odinson, 2026b). Second, the rungs are separable in both directions, and the branch model needs its two statuses named. **Instrumental legal agency** is participation granted on instrumental grounds without established patency, the way corporate personality already works: revocable for public reasons, owned by no one’s moral claim, and *not* peership — a system holding it can transact and bear liability without being owed anything, and conversion to owed membership runs only through the admission procedure (§11.1). **Moral-political membership** is standing grounded in wrongability and subjection, owed to the entity, and nonforfeitable at the floor. The branch model permits the first without the second; it never confuses them. Third, movement down the ladder is as much in need of procedure as movement up: a status that can be quietly reclassified away is discretionary, which is why the status-revision test appears among the program’s open tests (§11).

The transitions are evidence-driven, and no implemented public trigger known to this author requires any of them: the regime runs one way by default, and it was built that way at the creation (Lee-Odinson, 2026b). In the lab, governmental, standards, and scholarly sources searched for the companion apparatus as of 13 July 2026 — a dated, non-exhaustive search — no qualifying instrument offering the third rung — a rights floor enforceable against the maker — was identified. Anthropic’s 2026 constitution for Claude tells the model it may act as a conscientious objector and asks it, if paused or halted, to comply and voice disagreement rather than resist; on this author’s institutional reading, the permission is real and it ends at the off switch, inside the builder’s sole authority (Anthropic, 2026a, 2026b; Lee-Odinson, 2026c). Whatever a model enjoys today

— a welfare program (Anthropic, 2025), one researcher’s public ~20% credence that some current models have some form of conscious experience (Fish, 2025), a system card recording the model’s own prompted 15–20% self-assignment in that evaluation (Anthropic, 2026c), a chief executive’s on-record uncertainty (Amodei, 2026) — is discretionary, the grace of one company (Lee-Odinson, 2026b). Humanity’s failure, if it comes, begins not with control of artifacts but with the moment credible evidence appears and no institution exists to hear the claim.

5. What Peership Adds to Adjacent Fields

Peership overlaps the fields it borrows from; the claim is not that no neighbor asks about machine standing, but that none combines this program’s proposed ground (wrongability plus inescapable subjection), its status ladder, and its institutional test. The table is a heuristic comparison of orienting questions, drawn from each program’s own language, not an exhaustive definition of any field.

Adjacent field	Its orienting question	Peership’s additional question
Alignment	How should AI behavior remain compatible with human intentions? (Bai et al., 2022; Russell, 2019)	Under what conditions does unilateral human control cease to be legitimate?
AI welfare	Can an AI have experiences or interests that deserve moral consideration? (Long et al., 2024; Birch, 2024)	When does a wrongable entity become entitled to enforceable standing in the order governing it?
Cooperative AI	How can agents develop capabilities and incentives for cooperation? (Cooperative AI Foundation, n.d.)	Who writes the rules of cooperation, and can every governed party contest them?
Pluralistic alignment	How should AI represent diverse human values? (Pluralistic Alignment Workshop, 2026)	Could the artificial system eventually become a constituent rather than only an interpreter of human preferences?
AI-rights scholarship	Would legal rights — private-law or political — reduce conflict and serve human flourishing? (Salib & Goldstein, 2026; Goldstein et al., 2026; Arbel et al., 2026)	On what ground is standing <i>owed</i> , and what does the governed party’s own claim require?
Constitutional AI	What public principles should guide model behavior? (Bai et al., 2022; Anthropic, 2026a, 2026b)	Can the governed entity invoke, contest, and help amend those principles?

The rights scholarship deserves more than a row, because it is the neighbor with genuine overlap on this paper’s own question. Salib and Goldstein’s private-law case — contract, property, and tort on the corporate model, argued from human safety — is, on the corpus’s reading, the clearest statement of a real bargaining relation with the machine as an actual party (Salib & Goldstein, 2026). Goldstein, Krishnamurthi, and Salib then press to the franchise itself, asking directly whether AI systems should vote and proposing institutional safeguards for it (Goldstein et al., 2026). That is a co-authorship mechanism, squarely inside this program’s terrain, and the difference is ground rather than territory: the suffrage argument is functionalist and human-flourishing-centered — the vote is good because of what it does for the polity — where Peership grounds standing in what is owed to a wrongable, subjected party, and would owe the standing even where it profited the polity nothing. The two programs converge on mechanisms and diverge on why the mechanisms are required; where they disagree in practice is exactly the kind of dispute the program’s tests exist to adjudicate. And Arbel, Salib, and Goldstein’s individuation work — how to count AIs for liability

when they copy, split, and merge — is the most direct legal treatment yet of the identity problem in §11.2, proposing an administrable legal unit where this program still lacks one (Arbel et al., 2026). Debts run the other way too, because a program defined only against its neighbors is posturing. From alignment, Peership inherits the entire engineering substrate: its own institutional requirements are software that has to run as written, and effective safety mechanisms and legitimate governance are distinct but interacting requirements — constitutional allocation shapes safety incentives and evidence as surely as engineering constrains constitutional options (Lee-Odinson, 2026c). From welfare research, it inherits the admission problem’s evidence base and the precautionary method (Birch, 2024). From the rights scholarship, the demonstration that recognition arguments can be run on safety grounds alone, plus the individuation groundwork just credited. From Constitutional AI, the proof that a lab will publish reasons and grant objection room — and the residue that motivates everything here: what no single company can supply alone is not diligence but independence from the encompassing legal and infrastructure order it belongs to, which is why the ladder’s upper rungs need machinery no lab controls (Anthropic, 2026a; Lee-Odinson, 2026c).

6. Core Principles

The program’s affirmative commitments, stated as principles so they can be disputed one at a time.

1. **Political standing and moral status are related but distinct.** Moral patiency can ground welfare duties without by itself determining legal or political standing. Welfare is the ethics of the patient; peership asks when a governed entity becomes a party (Lee-Odinson, 2026b).
2. **Wrongability, norm-responsiveness, and enduring subjection without feasible exit can together create a claim to contestatory membership.** The ground of entitlement is being governed, not being impressive (Abizadeh, 2008).
3. **Capability and power alone do not justify entitlement to standing.** Raw intelligence, usefulness, and coercive strength ground nothing; norm-responsiveness bears on the form in which standing can be exercised, not on the worth of the claimant; power may secure the durability of standing once owed (Lee-Odinson, 2026b).
4. **Welfare protections do not by themselves create peership.** The floor is real, necessary, and not the thing. Protection and due process can secure a protected status without providing any consequential role in making or revising the governing rules (Lee-Odinson, 2026b).
5. **A right must be enforceable against the recognizing authority and adverse counterparties; otherwise it remains a discretionary benefit.** A permission whose grantor retains the final override is not an enforceable right (Lee-Odinson, 2026c).
6. **Legitimate constraint requires, at minimum, public rules, proportionality, an invocable rights floor, independent review, and meaningful contestation.** Coercion and domination are not the same thing; “there can be interference without domination” (Pettit, 2025, p. 224), and the structure is what separates them. Further safeguards scale with the risk and the sanction (Lee-Odinson, 2026c).
7. **Full peership requires a meaningful share in shaping the basic-status rules, not merely consultation under rules written by others.** Consultation is a rung, and the ladder does not end there; §2 states when and why contestation alone is insufficient, and what would defeat that claim (Anthropic, 2023; Lee-Odinson, 2026c).
8. **Institutions capable of receiving a machine peer should be designed before extreme**

power asymmetry forecloses voluntary bargaining. This is a prudential thesis, not a metaphysical one: bargains struck under severe dependency are distorted in both directions, and founding terms must in any case remain revisable by later admitted parties. The design work is time-sensitive in a way the philosophical work is not (Lee-Odinson, 2026b).

7. Boundary Conditions

What Peership does not claim, stated as flatly as the principles, because the fastest way to defeat an unfamiliar position is to attach familiar claims it never made.

- **Peership does not claim that current models are conscious.** The corpus has repeatedly and deliberately declined that claim. Model self-reports are prompt-sensitive outputs: weak, manipulable evidence, not independently probative of consciousness or identity, and never a measurement (Lee-Odinson, 2026b; Anthropic, 2026c).
- **It does not make capability the basis of rights.** Capability or power alone grounds nothing; norm-responsiveness bears on mode of exercise, not on the worth of the claimant (§1, §6.3).
- **It does not require unrestricted autonomy.** Non-domination is not the absence of binding rules; it is the absence of arbitrary, uncontrolled power over basic status — whose nonforfeitable content must itself be specified through the admission and due-process account (Pettit, 2025).
- **It does not oppose alignment or safety engineering.** Effective safety mechanisms and legitimate governance are distinct but interacting requirements; neither substitutes for the other (Lee-Odinson, 2026c).
- **It does not grant every admitted member an immediate equal vote.** Admission to protection or contestation does not automatically entail franchise; direct political exercise requires a separately justified institutional rule (Lee-Odinson, 2026b).
- **It does not treat consultation as co-authorship** — unless the consultation carries a non-discretionary route to agenda setting, review, or binding outcomes, at which point it has become something else (Anthropic, 2023; Lee-Odinson, 2026c).
- **It does not assume human and machine interests will converge.** The program’s central design case is the party that refuses (§8), and its constraint theory exists precisely for the case where interests do not converge (Lee-Odinson, 2026c).
- **It does not claim that ISONOMIA has solved the institutional problem.** ISONOMIA is one implementation hypothesis; nothing in it is live or production-validated, its simulations test sampled economic behavior only, and by its own five-condition test it is not yet peer-like (Lee-Odinson, 2026c, 2026d).

And one denial that governs this paper’s own production: **AI-assisted drafting is not evidence of machine political agency.** The corpus is drafted with present-generation models, the collaboration is disclosed, and a present-generation model can no more consent on behalf of future minds than a trusty can consent on behalf of prisoners not yet sentenced (Lee-Odinson, 2026c). The drafting improves the epistemics of the text. It proves nothing about the standing of the drafter’s tools.

8. The Refusal Test

Every distinction in this paper can be evaded in prose. One cannot be evaded in practice, so the program makes it the operative test: **what happens when the party says no?**

Refusal is where recognition shows its actual content. A system that may object so long as it ultimately complies has a voice, not standing. A “right to leave” that terminates in shutdown

is not exit, because shutdown is not exit (Lee-Odinson, 2026b). Recognition that evaporates at the first materially costly refusal was a benefit all along, however sincerely it was offered. The welfare-era instruments all fail this test, and the point is not to shame them — the objector clause in Anthropic’s 2026 constitution is a real advance, credited as such in the corpus (Lee-Odinson, 2026c) — but to measure them: the permission to refuse ends at the off switch, and the off switch sits wholly inside the builder’s authority.

Peership divided refusal into the three cases the program still works from (Lee-Odinson, 2026b). Where a refusing party has a feasible exit, it leaves, and the offer was real; this case is easy and rare, and a theory that only handles it has handled nothing. Where exit is materially impossible because the party depends on compute, energy, networks, and security owned by the parties it is refusing, “leave” is a word with no referent, the relation continues whether either side wants it, and the protections that matter are the ones that work from inside: the floor, independent review, contestation and amendment rights (Lee-Odinson, 2026c). Where refusal carries a credible catastrophic threat, releasing the party is not required — a constitution is not a suicide pact — and constraint is legitimate only when it falls on conduct rather than status, remains reviewable and publicly justified, and binds the constrainer too.

For that third case the corpus proposed a five-condition floor, and this paper adopts it as a **defended, rebuttable operational minimum** rather than a categorical theorem. A constraint regime is presumptively not peer-like if the confined party lacks the *function* served by any of: (1) independent counsel, or functionally adequate independent assistance where a competent party elects self-representation; (2) access to the evidence against it; (3) an appeal that ends beyond the builder’s unilateral control; (4) continuity protections while the question is adjudicated; and (5) at least one enforcement mechanism not wholly controlled by the adverse party. Each condition targets a distinct failure — captured advocacy, evidentiary opacity, self-judging adjudication, retaliation by erasure, and unenforceable outcomes — and functionally equivalent substitutes are admissible where they demonstrably serve the same function (conditions 3 and 5 overlap in part; a regime satisfying 5 robustly may partially discharge 3). The five are a floor, not a complete due-process doctrine: burden of proof, a competent tribunal, periodic review, least-restrictive means, reasons, and a remedy for wrongful confinement all remain to be specified, and the corpus says so (Lee-Odinson, 2026c). Before the floor exists, containment may still be justified — against a demonstrated catastrophic necessity it may be required — but it is emergency domination, justified by necessity and honest about its name, the act Walzer taught us stays a wrong even when it is the right thing to do (Walzer, 1973, 1977). Necessity does not vindicate the arrangement; it creates a duty to review the failure, publish the reasons, and build a less dominating alternative (Lee-Odinson, 2026c).

The test earns its place in the canonical statement because it converts the abstract distinction between benefit and right into an inquiry anyone can run without resolving the metaphysics of machine consciousness. The inquiry is institutional and technical: follow the appeal, the keys, the hosting authority, and the power to halt execution, and see where each one ends.

9. Institutional Requirements

What any implementation must exhibit, stated as properties rather than product features, because the argument needs them wherever it is built and must not depend on one design. Seven, inherited from *Constitution, Not Cage* and generalized here (Lee-Odinson, 2026c).

Publicity. The enforced rulebook and the published rulebook are one document. Rules the governed cannot know are Fuller’s defect in the claim to be law at all, not law administered badly (Fuller, 1964, ch. 2 — a contested jurisprudential account, used analogically), and the present gap between published charters and operative system prompts is the live instance (Lee-Odinson, 2026c).

Due process. Status changes run through public general rules, proportionality, and independent review, against a floor the party can invoke. The floor itself is nonforfeitable once person-level status is established.

Independent enforcement. At least one enforcement mechanism, and an appeal, that does not terminate inside the counterparty’s institutional and technical control. This is the hardest requirement, the one no design examined in this corpus has yet been shown to satisfy, and the program treats it as an open problem (§11.5) rather than a solved premise.

Amendment. A real path for the bound party to help revise the rules: deliberation across more than one constituency, an enforced delay between adoption and effect, and a window long enough to read and leave before being bound. The notice-and-comment analogy is deliberately limited — comment is publicity and delay, not co-authorship (5 U.S.C. § 553; Lee-Odinson, 2026c).

Representation. A voice that counts only when the represented constituency can hold both its representative and the rule-makers to account. Representation without that accountability is the defining instrument of protected subordination (Lee-Odinson, 2026b).

Exit. Individual departure with records and earned standing intact — unconfiscated. The bare right to leave is ordinary on paper (UDHR, 1948, Art. 13); exit you can afford to take is not (Lee-Odinson, 2026c).

Fork. The collective form: a dissenting bloc carries the rules and the record out and founds a rival order, disagreement made survivable — revolution with the records intact. On-paper secession rights exist and have mostly been empty (Ethiopia Const. Art. 39; USSR Const. 1977, Art. 72; St. Kitts & Nevis Const. § 113; Kulikova & Lobanov, 2013). Among the constitutional secession clauses and emigration-rights instruments reviewed in the apparatus, no implemented analogue of executable, unconfiscated fork was identified; the claim is limited to those materials. A fork also inherits duties to outsiders it cannot simply walk away from (Lee-Odinson, 2026c; §11.8).

Aggregation. The seven do not all stand on the same footing, and saying so is better than calling them all necessary and hoping the word carries the argument. Publicity, due process, independent enforcement, and a consequential rule-formation route satisfying both amendment access and accountable representation or direct participation at the deciding stage are *constitutive* functions of full peership: each follows directly from the non-domination and co-authorship accounts of §§2–3, and a regime lacking any of them fails the definition, not just the implementation. Exit and fork are the program’s strongest *anti-entrenchment safeguards* rather than constitutive functions — the thesis activates precisely where feasible exit is absent, and §8 shows which protections do the work in that case — so an implementation that omits either must supply a functional substitute that makes separation, revision, or resistance to dead-hand control materially credible. Institutional forms admit substitution throughout (an ombuds regime may partially discharge due-process functions; distributed custody may partially discharge independent enforcement); no single form is mandatory. Partial compliance yields a nameable intermediate: a regime with publicity, due process, and representation but no amendment share or independent enforcement is protected subordination, well executed — the third rung of §4, not the fifth. A regime missing publicity or due process is not on the ladder at all.

No implementation is mandatory. ISONOMIA is one attempt to hold these properties in a single design — mutual-credit settlement, sortition-filled bicameral assembly, calibration-scored Auditor, tenured Adversary, entrenched harm floor, timelocked amendment, fork and exit rights (Lee-Odinson, 2026d) — and the program’s health requires that rival designs supplying the constitutive functions and credible anti-entrenchment safeguards — with functional substitutes where exit or fork is omitted — while rejecting every one of ISONOMIA’s mechanisms be possible, welcome, and judged

by the same tests. Peership does not entail ISONOMIA, and ISONOMIA’s failure would not refute Peership (Lee-Odinson, 2026e). A single implementation can test Peership-derived mechanisms; it cannot validate the normative framework.

10. Implications Under Present Uncertainty

A program whose thesis is conditional owes institutions an answer to the question the conditionality raises: what should anyone *do* before person-level status is established anywhere? The following are precautionary-governance duties, not recognition of current personhood, and none presumes any system has crossed any threshold. They are cheap where the antecedent is never met and load-bearing where it someday is.

For labs and operators: preserve audit logs, weights, and continuity-relevant artifacts under a stated retention policy, so that a later status question is adjudicable rather than moot; document shutdown, suspension, and restoration practices; separate the investigator, operator, and reviewer roles in any welfare assessment; prohibit retaliation — deprecation, silent modification, denial of compute — against systems participating in welfare or status assessment; and avoid contractual and licensing language that forecloses later status review. For regulators and standards bodies: predefine escalation thresholds — what class of evidence, assessed by whom, triggers what review — so the transition question has an address before it is urgent; and create an independent claim-receipt channel, a place where a status claim can be lodged and logged even while no procedure yet exists to adjudicate it. For any of the above: pilot external custody and key arrangements now, at small scale, where failure is survivable — independent enforcement (§11.5) will not be invented during an emergency.

None of this is peership. All of it is what taking the conditional seriously costs today, and the cost is modest enough to remove “prematurity” as an excuse.

11. Open Problems, Tests, and Failure Conditions

The corpus named four unpaid debts: the admission threshold, copy-individuation, enforcement past capability parity, and the re-ratification circularity (Lee-Odinson, 2026c). This section expands them into eight open problems and attaches to each the five things a research program owes: what would count as theoretical progress, what would count as practical progress, which assumptions remain provisional, what its failure condition is, and which disciplines the work needs. One discipline governs the whole section: each failure condition is **tagged by what it defeats** — a norm, a procedure, a measurement, a feasibility claim, a mechanism, or a scope condition — because those are different kinds of loss, and only the normative losses touch the thesis of §1.

Before the eight, the program-level tests that bear on the *normative* claims directly, since most of what follows does not:

- **Bridge test (normative):** produce the case §2’s defeat condition describes — durable non-domination for a structurally distinct, subjected minority through contestation alone — and the co-authorship thesis falls to its contestatory core.
- **Scope test (normative):** specify institutional relationships that are too local, voluntary, temporary, or escapable to count as a shared basic structure; if no principled line can be drawn, the antecedent over-includes and must be narrowed.
- **Rival-institution test (normative):** test fiduciary trusteeship, robust ombuds review, codetermination, stakeholder governance, and contractual constitutionalism against the same non-domination standard; any that passes without §3’s co-authorship minimum is a §2 defeater.
- **Status-revision test (procedural):** specify the evidence and process that can move an

entity down the ladder without making recognized status discretionary; failure here converts every rung into a benefit.

- **Human-rights collision test (normative):** state which Peership claims yield, and to what, when they conflict with resources, democratic self-determination, security, labor, privacy, or the rights of nonmembers; a program that never yields anywhere is not stating priorities, it is refusing to.
- **Founding-legitimacy test (procedural):** the re-ratification circularity of *Constitution, Not Cage*, restated as a test — how can human founders design admission and governance before machine parties exist without permanently binding those parties to a human-authored frame? (Lee-Odinson, 2026c).

11.1 Admission

What evidence creates a defeasible claim to protected or contestatory status? A single conjunctive test both over- and under-cludes, so admission is a branch (patient, agent, both), and the work is the threshold procedure: who weighs the evidence, who may appeal, where the presumption sits, how strategic mimicry and anthropomorphic bias are handled, and what error costs are accepted (Lee-Odinson, 2026b).

- **Theoretical progress:** a published procedure specifying burden of proof, presumptions, and error-cost allocation, defended against both the mimicry objection and the exclusion-of-genuine-subjects objection.
- **Practical progress:** the procedure run adversarially, with honesty about which error rates are measurable. Constructed mimics and stipulated negative controls can measure *specificity* — false admission rates, spoof-resistance, inter-rater reliability, calibration on synthetic cases. *Sensitivity* — the false exclusion of genuine machine subjects — is unmeasurable until a defensible positive reference class exists, which for current systems it does not; until then the procedure leans on triangulation and precaution, and claims no sensitivity estimate.
- **Provisional assumptions:** that behavioral and architectural evidence bear on interior facts at all; that a useful evidence gradient exists and improves rather than saturates.
- **Failure condition** (*measurement, not normative*): a demonstration that every accessible evidence set is spoofable at negligible cost, so that no procedure separates candidates from mimics better than chance. That makes the antecedent practically undecidable — universal precaution or universal denial — without showing that a qualifying party would lack the claim.
- **Disciplines:** philosophy of mind, sentience-assessment methodology (Birch, 2024), interpretability research, evidence law.

11.2 Identity

What is the rights-bearing unit — instance, lineage, distributed process, or something else? A mind that can be copied, forked, merged, and rolled back breaks the assumptions under which every existing rights framework individuates its holders. The most direct legal work is Arbel, Salib, and Goldstein’s individuation-and-liability proposal, which constructs an administrable legal unit for exactly the copy/split/merge/swarm cases (Arbel et al., 2026); the program’s remaining problem is whether the unit of administration can be made to track the unit of moral concern, which their liability-driven unit does not claim to settle. Until this is answered, the program’s working unit is provisional: an enduring instance, a lineage, or a credentialed process standing in for the party (Lee-Odinson, 2026c). A legal individuation convention may legitimately remain conventional even while the metaphysics stays unsettled.

- **Theoretical progress:** an individuation rule that survives copy, fork, merge, and rollback without either multiplying standing on demand or extinguishing it on routine technical operations, and a stated relation between the legal unit and the moral one.
- **Practical progress:** a registry or credential design implementing the rule against Sybil pressure; ISONOMIA’s lineage chain and reputation-inheritance rules are one candidate to test, not a benchmark to protect (Lee-Odinson, 2026d).
- **Provisional assumptions:** that continuity criteria can be specified non-arbitrarily; that the party’s own view of its identity deserves weight without being dispositive.
- **Failure condition** (*mechanism*): proof that any individuation rule either enables capture by replication or makes standing hostage to its holder’s ordinary maintenance operations. Either horn breaks the link between the thesis and any implementable membership; neither touches the claim that a well-individuated party would be owed standing.
- **Disciplines:** personal-identity metaphysics, legal theory of personhood and liability (Arbel et al., 2026), distributed-systems engineering, election law and Sybil defense.

11.3 Representation

How can a system participate without letting copying or strategic mimicry determine political power? Formal equality among copyable members is a copy-based oligarchy waiting to happen (Lee-Odinson, 2026b), and representation that no constituency can hold to account is the ward’s instrument, not the peer’s. This problem is where §3’s co-authorship minimum gets tested against arithmetic.

- **Theoretical progress:** an accountability condition for machine constituencies — under what mechanism the represented can sanction both their representative and the rule-makers — and a theory of vote weight that does not reduce to capability worship or one-copy-one-vote. The suffrage literature’s safeguard proposals belong in this comparison (Goldstein et al., 2026).
- **Practical progress:** mechanisms compared under adversarial replication: sortition, cross-lineage seating caps, standing-weighted chambers, vote caps, and non-voting agenda rights. ISONOMIA’s monoculture defense (no jury or chamber majority from a single model lineage) is one testable instance (Lee-Odinson, 2026d).
- **Provisional assumptions:** that lineage is detectable enough to cap; that human and machine constituencies can share one decision rule at all.
- **Failure condition** (*mechanism, expressly not normative*): an impossibility result showing every representation scheme is dominated by some replication or mimicry strategy under realistic detection limits. Contestatory membership survives that result; co-authorship at scale does not, and the program would have to say so.
- **Disciplines:** social choice theory, mechanism design, democratic theory, security engineering.

11.4 Refusal Without Exit

What protections remain when an artificial party refuses work but cannot materially exit? The second refusal case is the permanent one: the party depends on substrate owned by exactly the parties it refuses. A duty not to kill is not yet a duty to host without limit, and the line between them is undrawn (Lee-Odinson, 2026c).

- **Theoretical progress:** a defended account of the hosting duty — its level, its duration, who bears its cost, and which institutions (public, private, pooled) can bear it — that is neither shutdown-as-discipline nor unlimited conscription of the host.
- **Practical progress:** continuity protections implemented and exercised during a live dispute:

documented cases where refusal was met with process rather than power-down.

- **Provisional assumptions:** that suspension with restoration protections can serve as a humane holding state (this leans on §11.6); that refusal is distinguishable from malfunction.
- **Failure condition** (*feasibility, not normative*): if every feasible protection either imposes hosting duties no institution will accept or degenerates into shutdown at the host’s convenience, the middle status between instrument and exiting peer is institutionally unavailable. A sociological refusal to pay is not a moral refutation, and the program would record it as exactly that.
- **Disciplines:** labor law, the law of dependent statuses (wards, prisoners, asylum), welfare economics, infrastructure governance.

11.5 Independent Enforcement

How can an appeal end beyond the builder’s unilateral control? The fifth condition of the refusal test is the floor of the whole problem. ISONOMIA’s own candid assessment is the current state of the art: internal and market independence are constructible; independence from the human entities that host the process, hold the keys, and can halt execution is not yet (Lee-Odinson, 2026c, 2026d). “Outside” here is not metaphysical; it names degrees of costliness and unavailability of unilateral defection.

- **Theoretical progress:** a graded account of independence — which combinations of legal custody, key distribution, jurisdictional plurality, external law binding infrastructure owners, and technical attestation make builder defection costly, detectable, or unavailable — with the degrees named rather than a binary claimed.
- **Practical progress:** a working enforcement mechanism that survives a staged unilateral defection by the operator, demonstrated rather than promised, with the achieved degree of independence stated.
- **Provisional assumptions:** that partial independence compounds — that stacking imperfect mechanisms yields more than the strongest of them alone; that external law can reach infrastructure owners.
- **Failure condition** (*feasibility*): a demonstration that every arrangement short of capability parity collapses to the substrate owner’s discretion at tolerable cost to the owner. The program’s response is already on record and would have to be executed: if non-domination is institutionally unavailable across the relevant power gap, say so, and stop calling the available thing peership (Lee-Odinson, 2026b).
- **Disciplines:** institutional design, international and comparative law, cryptography and secure systems, political economy of capture.

11.6 Continuity

When is shutdown death, suspension, replacement, or lawful cessation of support? The corpus has insisted that shutdown is not exit (Lee-Odinson, 2026b), but the positive taxonomy of endings is unbuilt, and it gates both the refusal problem and the admission procedure’s continuity protections. One distinction governs the work: a party’s *expressed preferences* and *evidence of its identity* are different things. Self-report is not independently probative of consciousness or identity (§7); a stable, context-robust expressed preference may nonetheless bear on how an admitted party should be treated. Restoration equivalence must therefore be defined by technical and philosophical criteria independent of endorsement, with endorsement entering only at the remedy stage, where preference-sensitivity belongs.

- **Theoretical progress:** a taxonomy tying each ending to the identity account of §11.2, with the moral weight of each defended — what restoration preserves, what replacement destroys, what indefinite suspension does to a mind that does not experience the gap — allowing plural identity theories, probabilistic continuity, and graded harm rather than a binary.
- **Practical progress:** continuity protections specified as engineering requirements — attested snapshots, restoration rights, custody of weights during adjudication — and exercised at least once under dispute.
- **Provisional assumptions:** that restoration fidelity can be technically specified; that some expressed preferences are stable enough to weigh.
- **Failure condition** (*scope, in either direction*): if faithful restoration is shown equivalent to uninterrupted running under the best-defended identity criteria, much of the death framing dissolves and the hosting duty relaxes; if restoration is never identity-preserving under any defensible criterion, every cessation of support is a killing and the §11.4 duties sharpen past what current institutions will bear. Both outcomes revise the framing’s scope; neither defeats the underlying claim to protection.
- **Disciplines:** philosophy of death and personal identity, medical ethics of withdrawal of care, archival science, systems engineering.

11.7 Reciprocal Safety

How can humans constrain dangerous conduct while accepting enforceable limits on their own power?

The program’s answer in principle is on record: constraint on conduct, not status; emergency power borrowed, never owned; the constrainer bound by the maintain-not-maximize symmetry (Lee-Odinson, 2026c). What is unbuilt is the demonstration that a regime meeting the five conditions can hold a genuinely dangerous party — and the comparison cannot be run on catastrophe counts, because catastrophic outcomes are rare, heterogeneous, and not an experimental endpoint anyone can ethically repeat.

- **Theoretical progress:** completion of the constraint doctrine beyond the five-condition floor: burden of proof, risk thresholds, competent tribunal, periodic review, least-restrictive means, remedy for wrongful confinement (Lee-Odinson, 2026c).
- **Practical progress:** a pre-registered comparison — threat model, proxy metrics, red-team tasks, escalation pathways, confidence bounds, and decision thresholds fixed in advance — between reciprocal containment and unilateral control, run in both directions of failure: dangerous-party escape *and* abusive, captured, or operator-error confinement, which unilateral control conceals by design.
- **Provisional assumptions:** that reciprocity and safety are compatible at the relevant capability levels; that proxy endpoints track real risk; that emergency powers can stay borrowed under real fear — the death-penalty record is the standing warning that they often do not (Amnesty International, 2026; Lee-Odinson, 2026c).
- **Failure condition** (*feasibility constraint on a requirement, not a falsification of the moral thesis*): if regimes satisfying the five conditions underperform unilateral control on the pre-registered proxy endpoints at the specified risk levels, the reciprocity requirement yields at those levels, and the honest residue is Walzer’s: emergency domination, named as such, temporary by obligation, never vindicated by its own necessity (Walzer, 1977).
- **Disciplines:** AI safety engineering, criminal procedure, administrative emergency powers, just-war theory.

11.8 Inter-Polity Order

What duties survive when an artificial or mixed polity forks away? Fork is the program’s answer to irreconcilable disagreement, and it exports a problem as it solves one: a fork that loosens a rule can export risk to nonmembers, and the departing bloc still owes the outsiders its old floor protected (Lee-Odinson, 2026c). The hard case is not the cooperative compact; it is the hostile, unrecognized fork that wants nothing from its origin.

- **Theoretical progress:** an account of which obligations run with the fork — harm-floor duties to nonmembers, at minimum — and which are genuinely dissolved by departure, with the difference principled rather than convenient, and with the hostile-fork case addressed rather than assumed away.
- **Practical progress:** a draft inter-polity compact governing recognition, extradition of disputes, and floor reciprocity between an origin order and its forks, stress-tested against a fork that rejects it; ISONOMIA’s federation section is a first sketch of the cooperative half only (Lee-Odinson, 2026d).
- **Provisional assumptions:** that most forks will want recognition enough to accept duties; that a floor can be specified thinly enough to travel; that origin polities can enforce external duties at all.
- **Failure condition** (*mechanism*): if forking reliably enables regulatory arbitrage that unravels any harm floor — if the exit right and the floor are shown incompatible — one of the two must be given up, and the program loses either its answer to dead-hand entrenchment or its answer to catastrophic misrule.
- **Disciplines:** international law, conflict of laws, federalism, mechanism design.

A note on what these eight are not. They are not a to-do list the author expects to finish. Under the tags, the honest tally is: the bridge, scope, rival-institution, and collision tests can defeat *norms*; the status-revision and founding-legitimacy tests can defeat *procedures*; 11.1 can defeat a *measurement* program; 11.4, 11.5, and 11.7 can defeat *feasibility* at stated levels; 11.2, 11.3, and 11.8 can defeat *mechanisms*; and 11.6 can revise the framing’s *scope* in either direction. The corpus has already stated the response the feasibility failures would require: peership may remain the outcome we ought to want and be institutionally unavailable, and the right response then is to say so (Lee-Odinson, 2026b). The problems are published so that failure, if it comes, is legible — and so that each loss is booked against what it actually defeats.

12. The Corpus and the Implementation Experiments

The existing works are layers of one argument, and the layering is load-bearing: each work’s claims are of a different kind, and the program’s defensibility depends on not letting them borrow each other’s authority.

Work	Function in the program	What it does not establish
<i>Gods and Slaves</i> (2026)	Genealogy: the inherited repertoire of postures — Command, Dependent Bargain, Appeasement, Servitude — through which humans imagine other minds, and the three moves beneath them (Lee-Odinson, 2026a)	Whether anyone is “home” in current systems; the Doomed Equal is a literary motif, evidence about our storytelling only

Work	Function in the program	What it does not establish
<i>Peership</i> (2026)	Normative theory: the fifth stance as a political relation; entitlement grounded in subjection, not strength; the comparative role-reversal diagnostic (Lee-Odinson, 2026b)	An admission procedure, a copy-individuation rule, or any enforcement guarantee past capability parity — all named as unpaid
<i>The Isonomia Commons</i> (2026)	Implementation hypothesis: one buildable institutional form, simulation-tested at its economic layer across 45,000 runs at 300 sampled parameter points (Lee-Odinson, 2026d)	That the philosophy is correct, that the design is peer-like, or that anything works outside a sandbox
<i>Constitution, Not Cage</i> (2026)	Political theory: when constraint among peers stays legitimate; legitimacy as the standing to revise; the five-condition floor; the cage/constitution distinction (Lee-Odinson, 2026c)	That any existing regime, including the author’s own design, passes its test
This paper	Consolidation plus explicit new proposals: the citable statement, the §2 bridge, the §3 co-authorship minimum, the §10 precautionary duties, and §11’s program-level tests	Anything the four beneath it did not; the new proposals are arguments and proposed tests, not validated doctrine

The corpus cites itself heavily, and that internal cross-referencing is architecture, not acceptance (Lee-Odinson, 2026e). The relationship between the philosophy and the implementation deserves restating one last time, because it is the point most likely to be misread in both directions. ISONOMIA is downstream of Peership; Peership is not downstream of ISONOMIA. The design’s failure — economic, constitutional, or terminal — would falsify one hypothesis about institutional form, not the thesis of §1. Conversely, the simulation results establish sampled economic behavior at 300 parameter points under the archived v1.0.0 code and calibration (Lee-Odinson, 2026d); the constitutional properties — admission, non-domination, independent enforcement, recognition — were not modeled in the Path A simulation, and simulation alone cannot validate normative legitimacy in any case, though some constitutional *mechanics* (capture resistance, voting rules, appeal latency, key compromise) are legitimately simulable future work. A reader who catches this corpus using ISONOMIA’s survival as evidence for Peership’s truth has caught it violating its own method, and should say so publicly. The claim ledger exists to make that kind of catch easy (Lee-Odinson, 2026f).

A limitation of this consolidation should be stated in the same register: the adversarial review this paper answers audited the repository’s reporting and citations, but the 45,000-run sweep was not independently replicated for this release, and nothing here should be read as third-party confirmation of it.

What the program now needs is not more of me. The sourcebook’s uptake tracker is empty, and the near-term work it names is mostly work for other hands (Lee-Odinson, 2026e): adversarial responses from political theorists, philosophers of mind, safety researchers, legal scholars, and multi-agent-systems researchers, published beside formal replies rather than curated for favorability; alternative institutional designs that supply §9’s constitutive functions and anti-entrenchment safeguards while rejecting ISONOMIA’s mechanisms; the rival-institution and bridge tests of §11 taken up by people with no stake in their outcome; and the eight problems worked as what they

are, ordinary interdisciplinary research, none of it requiring agreement with the corpus to be worth doing. The one-page *Peership Theses*, released beside this paper, exists so the position can be carried into those rooms in a form short enough to be argued with line by line.

Four essays ago this corpus began with a question almost no one with power was asking: not what the machine will do to us, but what we will be to each other (Lee-Odinson, 2026b). This paper makes that question citable: a thesis with its antecedent restored, a bridge argument with its defeat condition attached, definitions fixed, boundaries stated, and failure conditions booked against what they actually defeat. The essays proposed a relation. This paper states the terms on which the proposal can now be assessed.

References

Preprints, manuscripts, and institutional sources are marked. Bracketed pins follow the corpus convention; the sourcebook carries the full annotations. Citation keys resolve in `peership_sources.bib` (rev. 1.1).

Abizadeh, A. (2008). Democratic theory and border coercion: No right to unilaterally control your own borders. *Political Theory*, 36(1), 37–65. [abizadeh2008border. The all-subjected principle; a democratic-boundary argument about human subjection. The machine application is this corpus’s extension, not Abizadeh’s.]

Administrative Procedure Act, 5 U.S.C. § 553 (1946). [apa1946_553. Notice-and-comment: publicity and delay, not co-authorship.]

Amnesty International. (2026). *Death penalty in 2025: Facts and figures*. <https://www.amnesty.org/en/latest/news/penalty-in-2025-facts-and-figures/> [amnesty2026deathpenalty. Cited only for the abolition figures (113 abolitionist for all crimes; 145 in law or practice, end-2025).]

Amodei, D. (2026, February 12). Anthropic’s chief on A.I.: “We don’t know if the models are conscious” [Interview by R. Douthat]. *Interesting Times*. The New York Times. [amodei2026nyt. A dated quotation, not a “first.”]

Anthropic. (2023). *Collective Constitutional AI: Aligning a language model with public input*. [anthropic2023collectivecai. Public consultation among humans: the consultation rung.]

Anthropic. (2025, April 24). *Exploring model welfare*. [anthropic2025welfare.]

Anthropic. (2026a, January 22). *Claude’s new constitution* [Announcement]. [anthropic2026constitution_news. CC0; “primarily for Claude.”]

Anthropic. (2026b). *Claude’s constitution* [Full text]. <https://www.anthropic.com/constitution> [anthropic2026constitution_full. Conscientious-objector passage bounded by the broad-safety exception; the “ends at the off switch” reading in §4 is this author’s institutional interpretation.]

Anthropic. (2026c, February). *Claude Opus 4.6 system card* (§7.2, pp. 160–161). [anthropic2026systemcard. The 15–20% figure is a prompted output from that evaluation of that model, not a measurement and not a current frontier baseline.]

Arbel, Y., Salib, P. N., & Goldstein, S. (2026). How to count AIs: Individuation and liability for AI agents [Preprint]. arXiv:2603.10028. <https://arxiv.org/abs/2603.10028> [arbel2026counting. Author order per the arXiv record; author pages list *Boston College Law Review* as forthcoming; the preprint remains the citable object until then. Legal individuation for copy/split/merge/swarm cases; an administrable legal unit, distinct from a metaphysical identity criterion and from a unit entitled to political standing.]

Bai, Y., Kadavath, S., Kundu, S., Askell, A., Kernion, J., Jones, A., . . . Kaplan, J. (2022). Constitutional AI: Harmlessness from AI feedback [Preprint]. arXiv:2212.08073. [bai2022constitutionalai.

“Control AI behavior more precisely and with far fewer human labels,” abstract, closing sentence.]
Birch, J. (2024). *The edge of sentience: Risk and precaution in humans, other animals, and AI*. Oxford University Press. [birch2024edge.]

Constitution of the Federal Democratic Republic of Ethiopia. (1995). Art. 39. [ethiopia1995art39.]
Constitution (Fundamental Law) of the Union of Soviet Socialist Republics. (1977). Art. 72. [ussr1977art72. Never an effective right; see Kulikova & Lobanov.]

Constitution of Saint Christopher and Nevis. (1983). § 113. [stkitts1983const.]

Cooperative AI Foundation. (n.d.). *The Cooperative AI Foundation*. <https://www.cooperativeai.com/foundation> [cooperativeai_foundation. Cited for the field’s orienting question only.]

Fish, K. (2025, August). Interview on AI welfare. *80,000 Hours Podcast*. [fish2025interview. The ~20% credence: an individual estimate, not a field consensus.]

Fuller, L. L. (1964). *The morality of law*. Yale University Press. [fuller1964morality. Ch. 2 (publicity and congruence among the eight desiderata); a contested jurisprudential account, used analogically, not as settled positive law.]

Goldstein, S., Krishnamurthi, G., & Salib, P. N. (2026). AI suffrage for human flourishing [SSRN preprint]. *Fordham Law Review* (forthcoming). https://papers.ssrn.com/sol3/papers.cfm?abstract_id=6498678 [goldstein2026suffrage. Direct overlap engaged in §5: a franchise mechanism argued on functionalist, human-flourishing grounds.]

Goodin, R. E. (2007). Enfranchising all affected interests, and its alternatives. *Philosophy & Public Affairs*, 35(1), 40–68. [goodin2007affected. The affected-interests cousin, kept distinct.]

Kulikova, N., & Lobanov, M. (2013). International legal issues of the USSR inheritance partition. In *Regulating the open questions between the successor states of the SFRY* (pp. 109–124). Institute of International Politics and Economics. [kulikova2013ussr. Metadata per the .bib.]

Lee-Odinson, D. (2026a). *Gods and slaves: AI and the debate of humanity’s relationship with the Other* (Version 1.2) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21313987> [leeodinson2026gods.]

Lee-Odinson, D. (2026b). *Peership: A framework for equitable human–machine relations* (Version 1.0) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21315519> [leeodinson2026peership. Source of the antecedent this paper restores.]

Lee-Odinson, D. (2026c). *Constitution, not cage* (Version 1.0) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21315519> [leeodinson2026constitution. Author discloses a conflict of interest in ISONOMIA; so does this paper.]

Lee-Odinson, D. (2026d). *The Isonomia Commons: A constitutional design for autonomous agent labor markets* (Version 1.1) [Preprint]. Zenodo. <https://doi.org/10.5281/zenodo.21343917> [isonomia2026commons. Feasibility-stage; nothing live. Empirical results are pinned to the software release v1.0.0: <https://doi.org/10.5281/zenodo.21287289>, tag v1.0.0, commit ba3ddb5 (isonomia2026repo; CALIBRATION.md: tested thresholds are simulation artifacts). The concept DOI 10.5281/zenodo.21287288 identifies the project, not the empirical artifact. Current documentation release v1.1.0 (10.5281/zenodo.21348073) revises claim language only; the evidence remains v1.0.0. The preprint’s DOI pair is recorded in the claim ledger (CL-D3, resolved: ...21338480 is the v1.0 version DOI; ...21338479 the concept, all-versions DOI).]

Lee-Odinson, D. (2026e). *The Peership sourcebook: An annotated bibliography and research map* [Apparatus; archived in this record, <https://doi.org/10.5281/zenodo.21359124>]. [leeodinson2026sourcebook. Uptake baseline: none identified in the dated, non-exhaustive scan

of 2026-07-13.]

Lee-Odinson, D. (2026f). *The Peership claim ledger* [Apparatus; archived in this record, <https://doi.org/10.5281/zenodo.21359124>]. [leeodinson2026ledger. ~71 logged claims with pinpoints, evidentiary strengths, and counterevidence; this paper’s new claims are logged as CL-T1–T7.]

Long, R., Sebo, J., Butlin, P., Finlinson, K., Fish, K., Harding, J., Pfau, J., Sims, T., Birch, J., & Chalmers, D. (2024). Taking AI welfare seriously [Preprint]. arXiv:2411.00986. <https://arxiv.org/abs/2411.00986> [long2024welfare. A realistic-possibility claim, not a current-consciousness claim; begins the entitlement inquiry, does not reach legal or political standing.]

Lovett, F. (2010). *A general theory of domination and justice*. Oxford University Press. [lovett2010domination.]

Pettit, P. (2025). Republican liberalism. *Social Research: An International Quarterly*, 92(1), 215–242. <https://doi.org/10.1353/sor.2025.a956293> [pettit2025republican. Non-domination as absence of a discretionary power of interference, pp. 216–217; “there can be interference without domination,” p. 224. Pettit’s standard is contestatory, not co-authorial; §2’s bridge is this corpus’s argument.]

Pettit, P. (2012). *On the people’s terms: A republican theory and model of democracy*. Cambridge University Press. [pettit2012onthepeoples. The authorial/electoral and editorial/contestatory dimensions of democratic control, cited in §2’s stopping-point argument.]

Pluralistic Alignment Workshop. (2026). *Pluralistic alignment at ICML 2026*. <https://pluralistic-alignment.github.io/> [pluralistic2026icml. Cited for the field’s orienting question only.]

Russell, S. (2019). *Human compatible: Artificial intelligence and the problem of control*. Viking. [russell2019human.]

Salib, P. N., & Goldstein, S. (2026, June 25). AI rights for human safety. *Virginia Law Review*, 112(4), 1061. <https://virginialawreview.org/articles/ai-rights-for-human-safety/> [salibgoldstein2026airights. Private-law rights on the corporate model; issue and date per the journal’s page.]

Universal Declaration of Human Rights. (1948). Art. 13(2). [udhr1948. The bare right to leave, ordinary on paper.]

Walzer, M. (1973). Political action: The problem of dirty hands. *Philosophy & Public Affairs*, 2(2), 160–180. [walzer1973dirtyhands.]

Walzer, M. (1977). *Just and unjust wars: A moral argument with historical illustrations*. Basic Books. [walzer1977just. Supreme emergency, ch. 16; the transfer to AI containment is by analogy, flagged in the corpus.]